At the Edge of AI at the Edge
Ultra efficient AI on low-power devices

Mohammad Rastegari
mohammad@xnor.ai

May 2018

XNOR.AI

AI is confined to the cloud

far from the users at the edge

Goal

XNOR.AI bridges the growing divide between AI models dependent on the cloud and devices running at the edge

Deep learning models reliant on the cloud

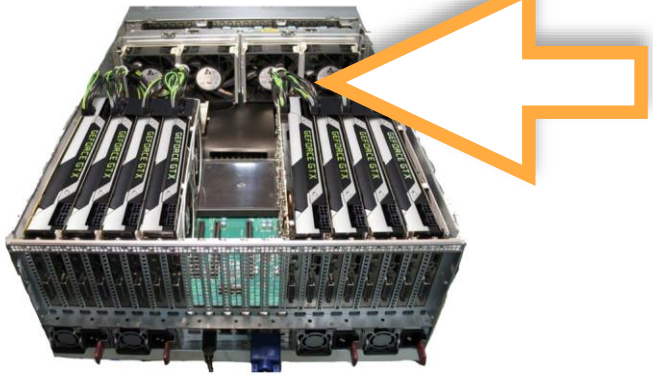

Growing demand for edge devices

Opportunities

Intelligent cameras that preserve privacy, security and bandwidth at home

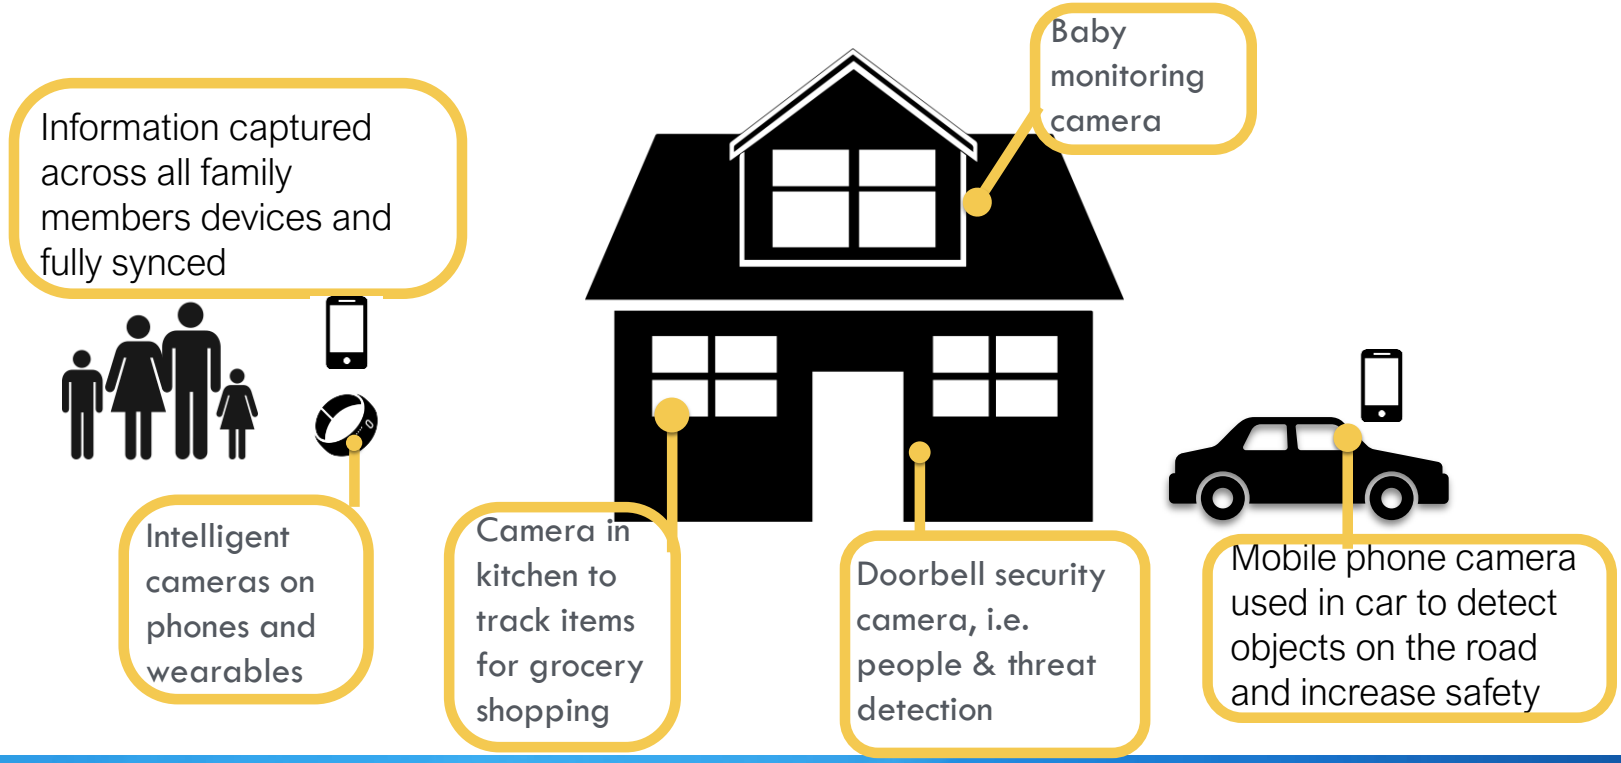

Opportunities

Intelligent cameras for Advanced Driver Assistance Systems (ADAS)

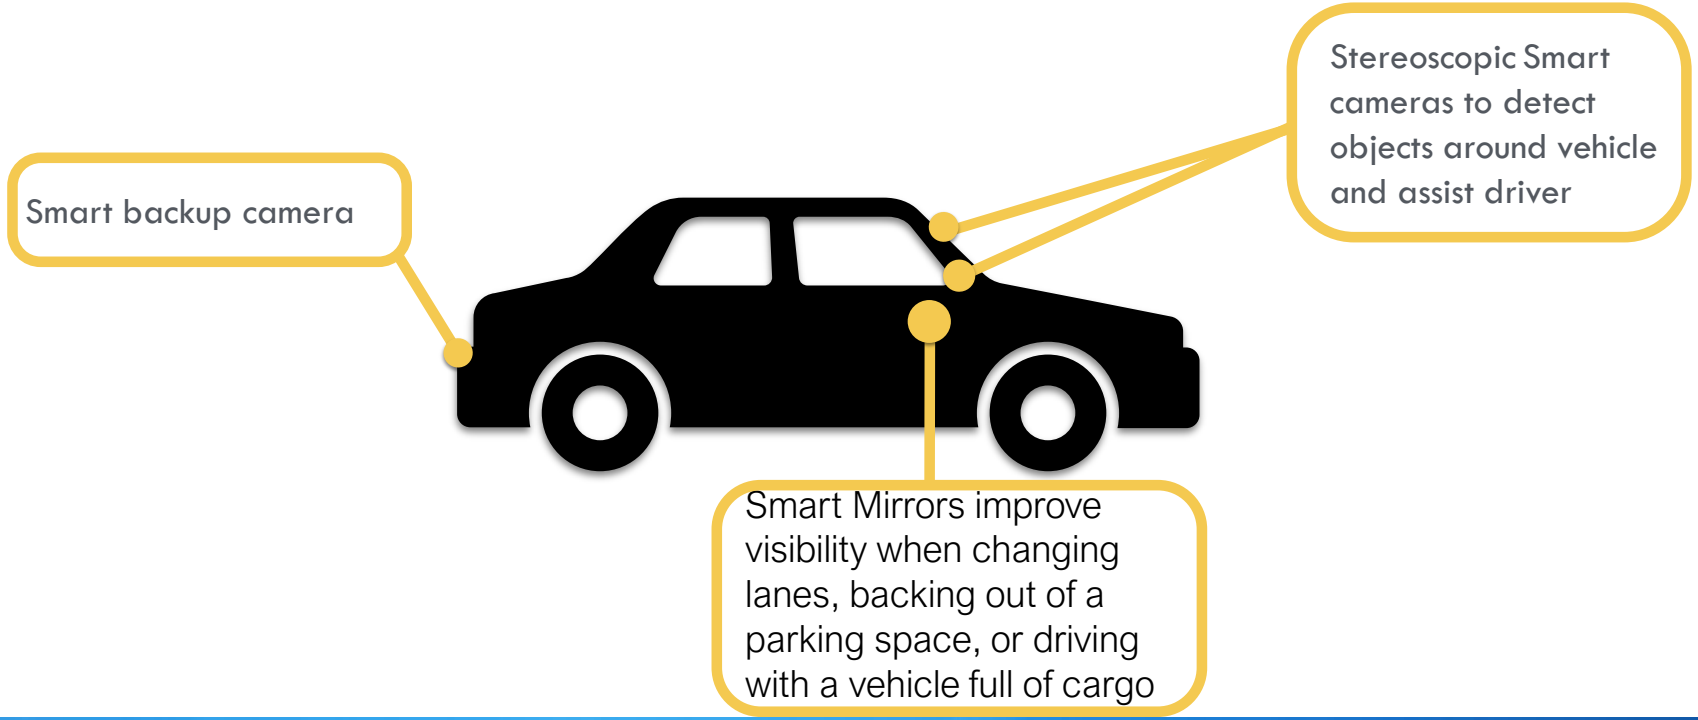

Opportunities

Intelligent cameras provide real-time tracking, recognition & detection on device

Automate remote monitoring

Intelligent wearable cameras provide personal security

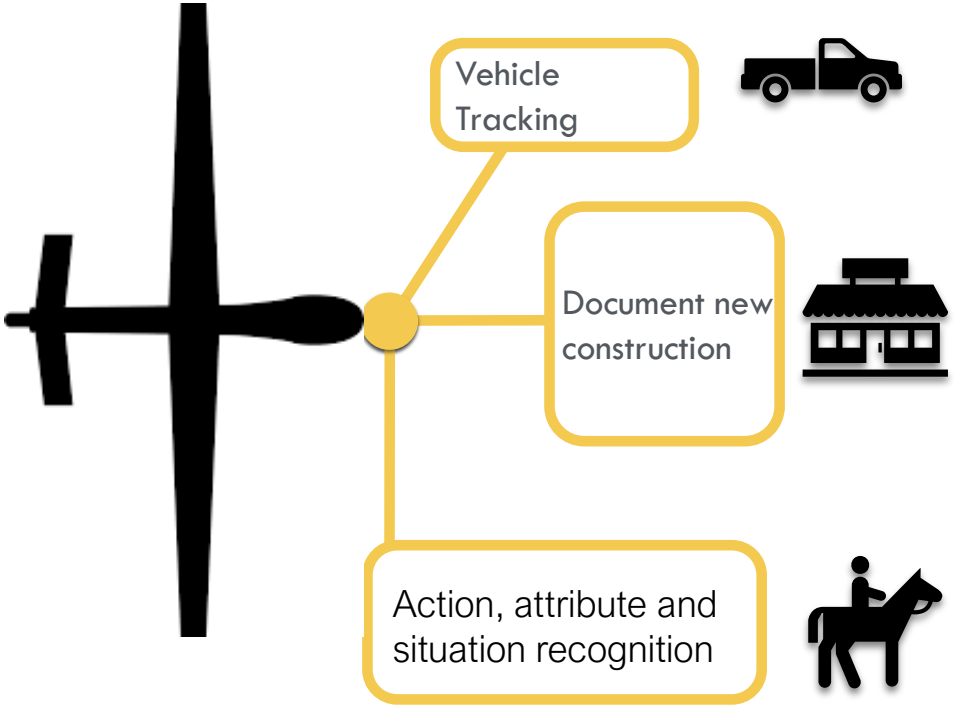

What we do

Image Tagging

Image Enhancement

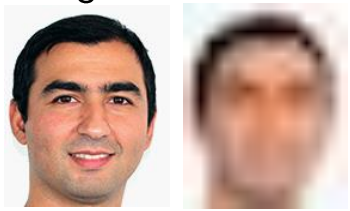

Tracking

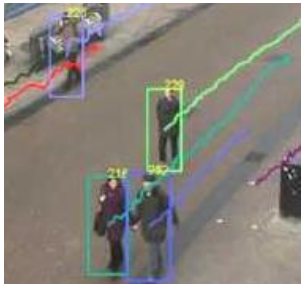

Action Recognition

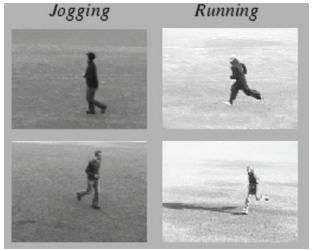

Object Detection

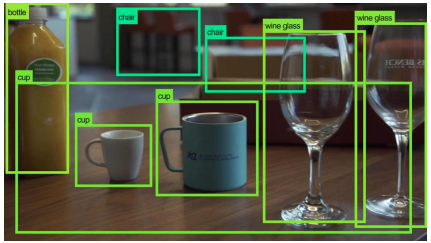

Scene Recognition

Segmentation

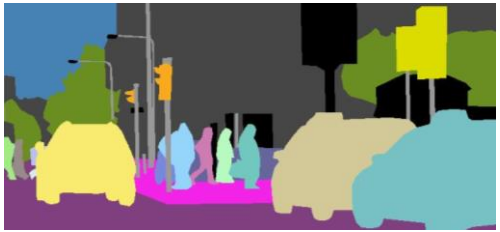

Object Detection: An Expensive Task in AI

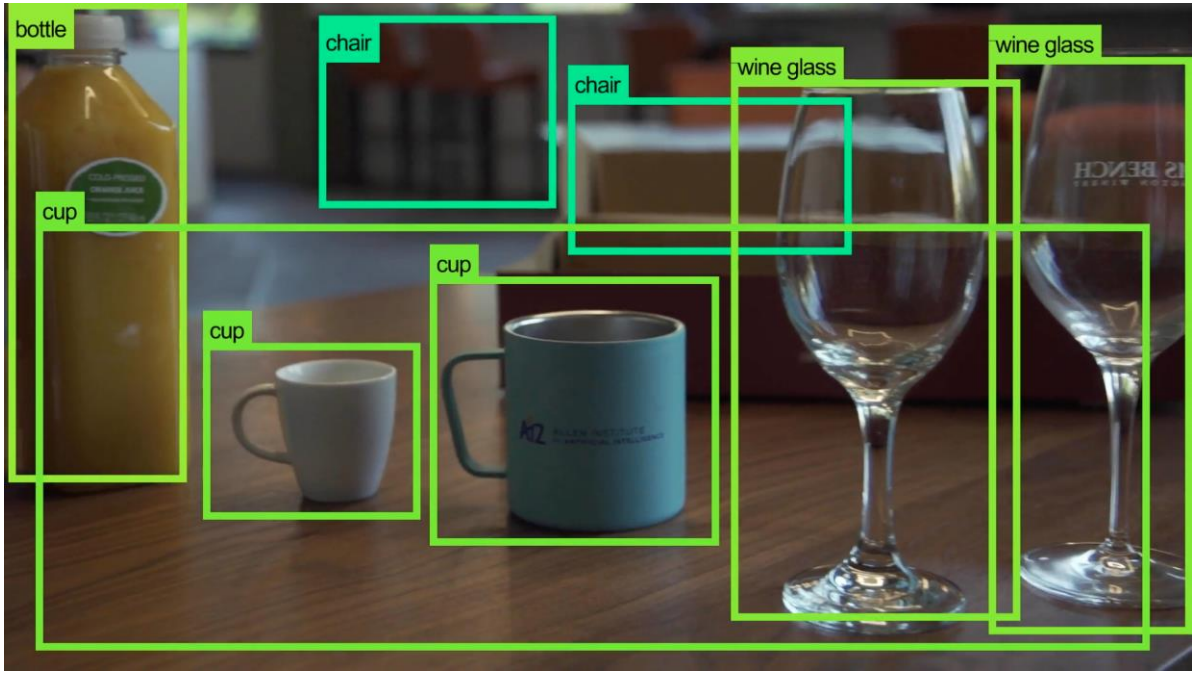

Dash Cam

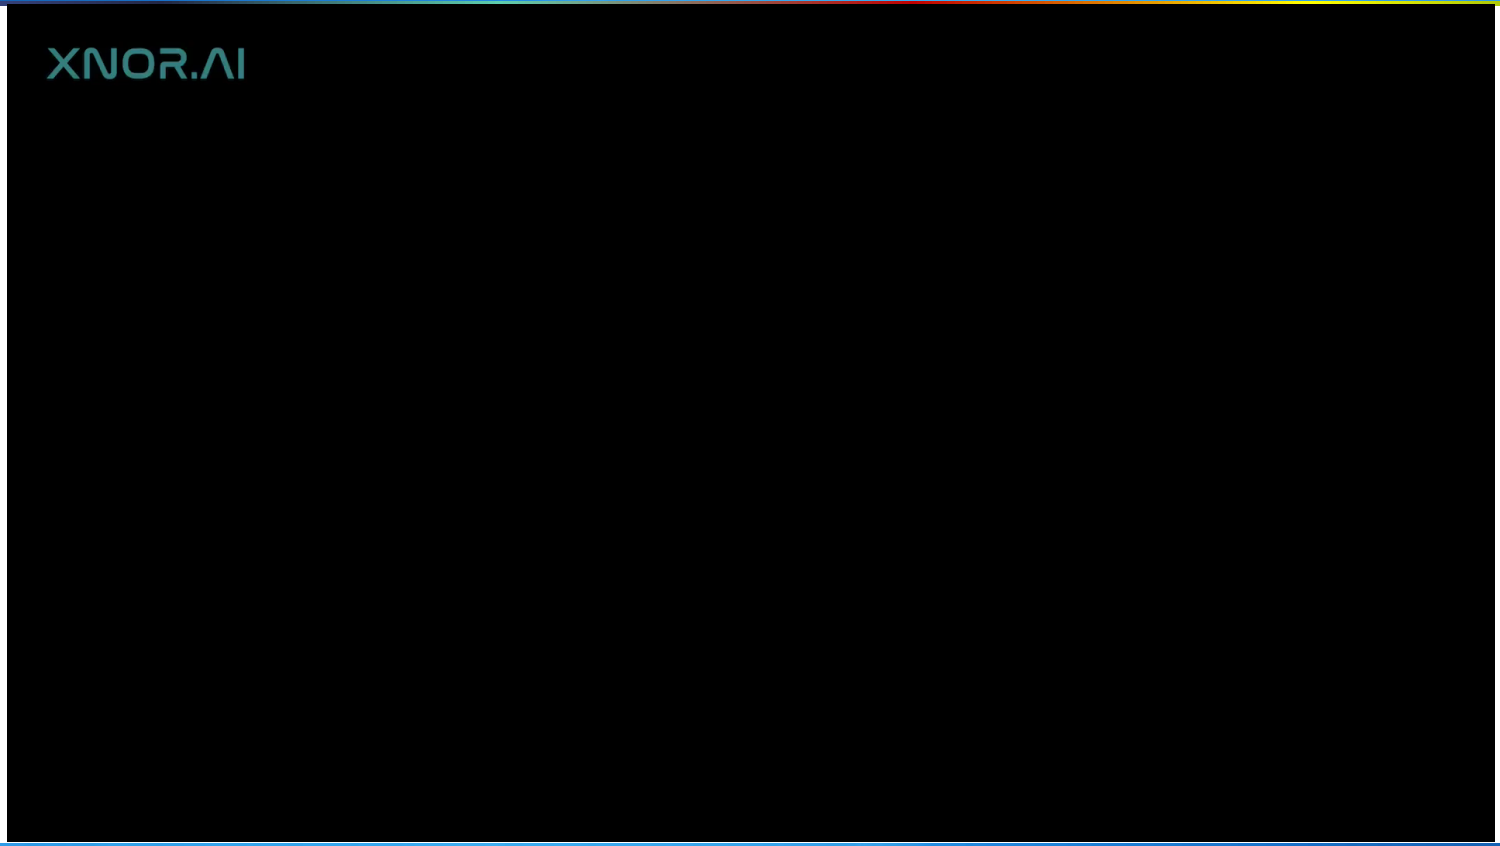

Security Camera

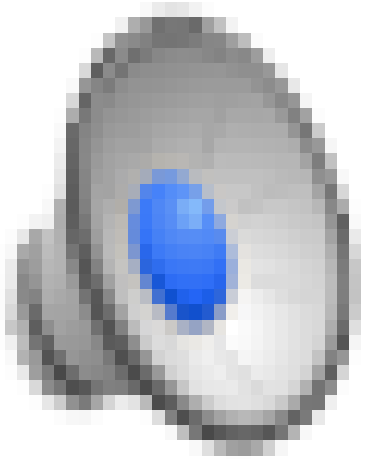

Fine Grain Categories

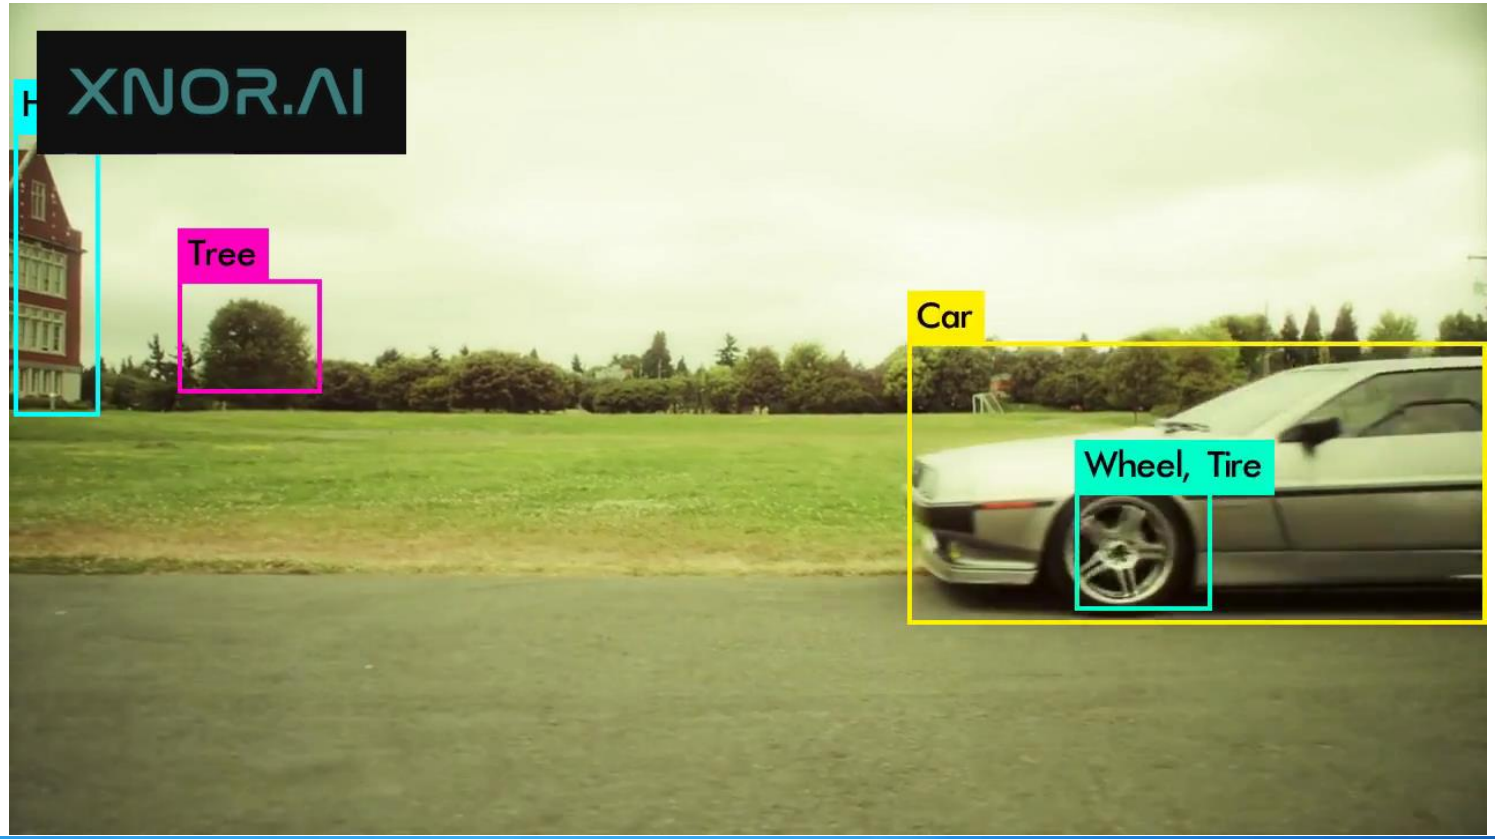

Hardware Platforms

XNOR.AI provides fast and efficient AI at the edge

Low Power AI

State-of-the-Art AI: all the way to Pi Zero

XNOR $5 deep learning machine…
… on Raspberry Pi Zero

Modular AI at Edge

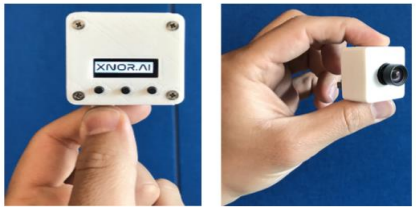

Low Power AI

Ambarella S5L

- Very low power (~2x lower than Pi Zero)
- Standard AI model for object detection
  - 1 fps
- XNOR AI Model for object detection
  - 17 fps

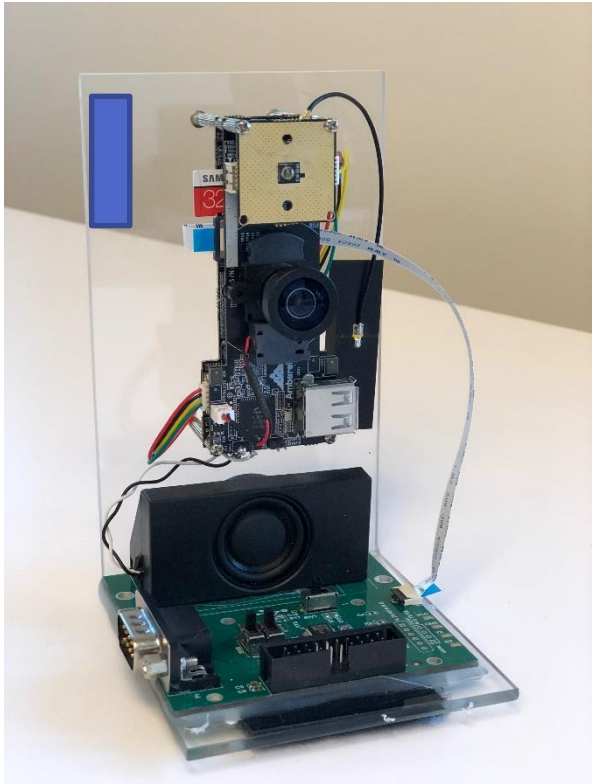

How our technology works

Convolutional Neural Network

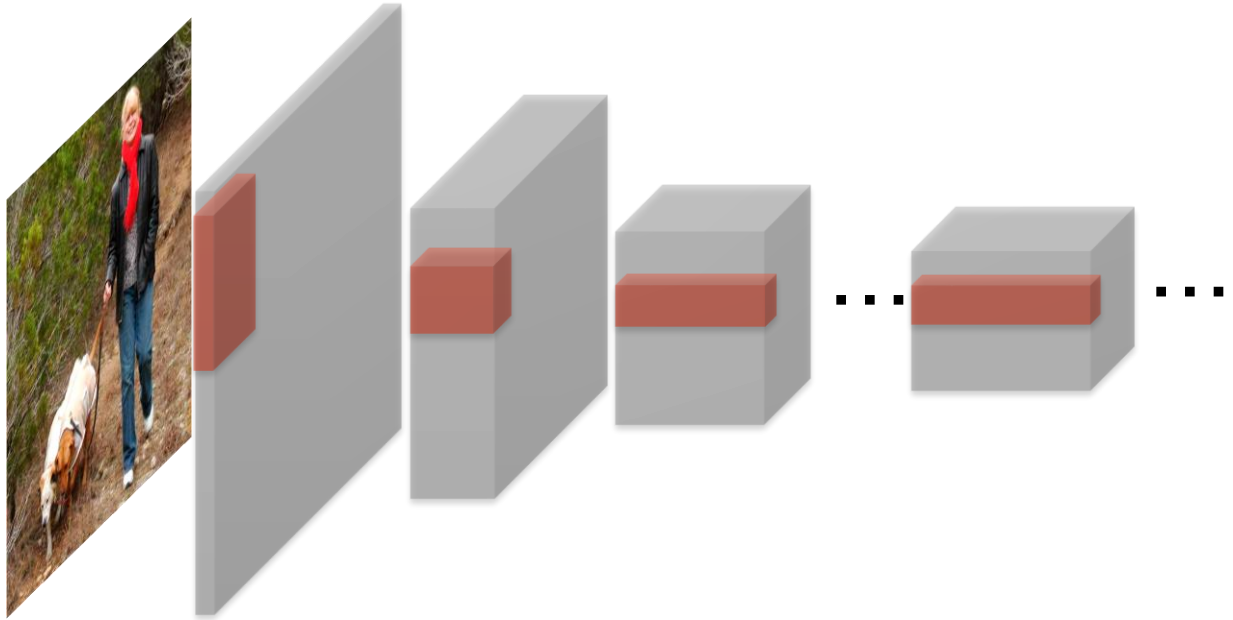

GPU !

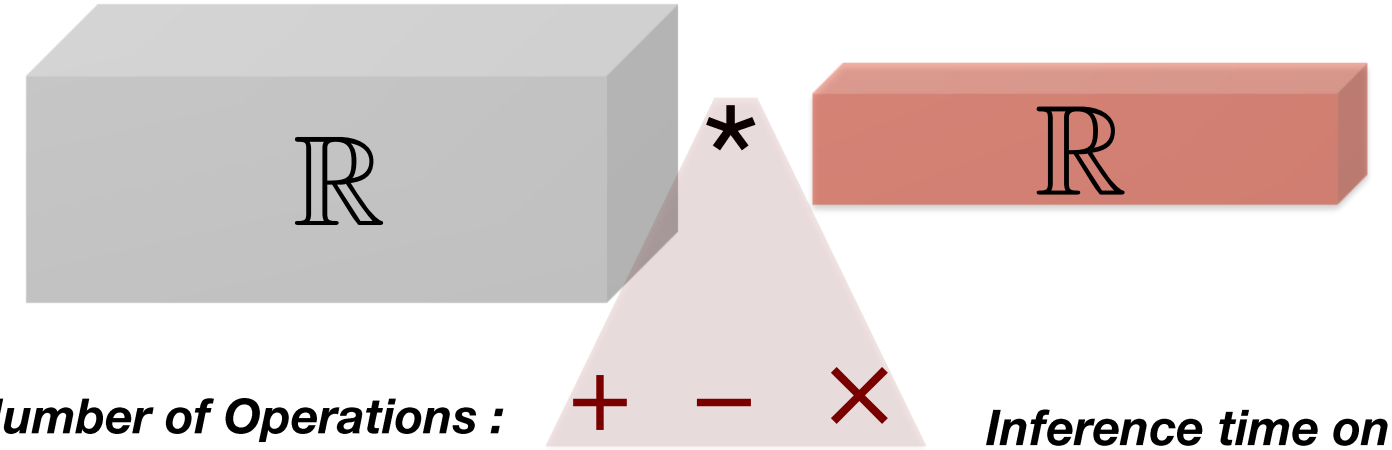

Number of Operations :

- AlexNet →1.5B FLOPs
- VGG → 19.6B FLOPs

Inference time on CPU :

- AlexNet →~3 fps
- VGG → ~0.25 fps

Lower Precision

Reducing Precision
- Saving Memory
- Saving Computation

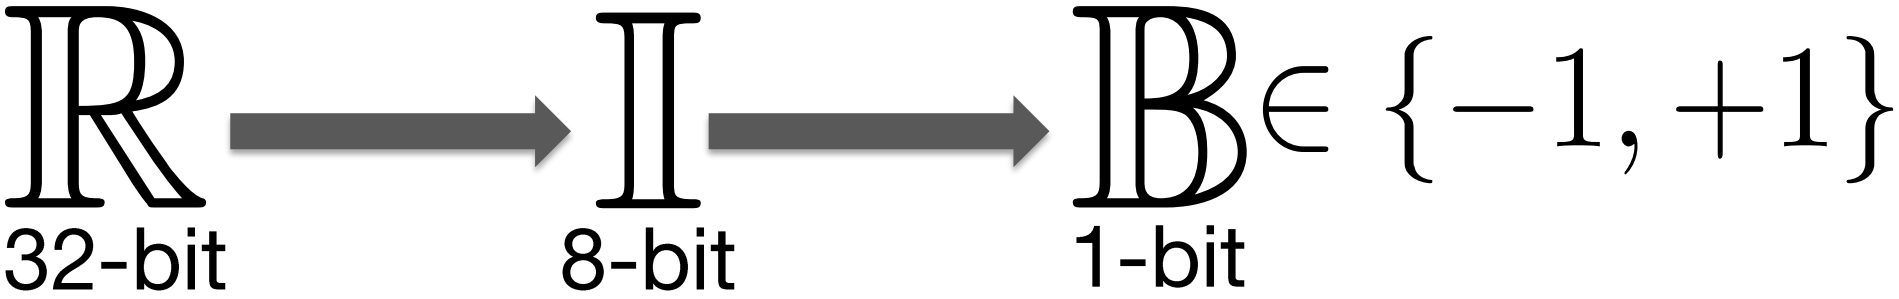

{-1,+1}	{0,1}
MUL	XNOR
ADD, SUB	Bit-Count (popcount)

Why Binary?

- Binary Instructions
  - AND, OR, XOR, XNOR, PopCount (Bit-Count)

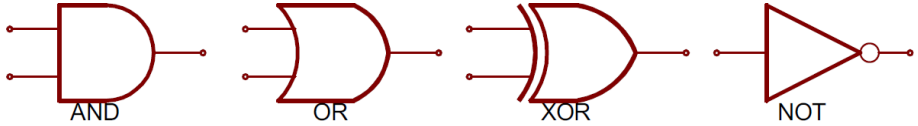

- Low Power Device

- Easy to Implement in hardware

XNOR.AI Technology

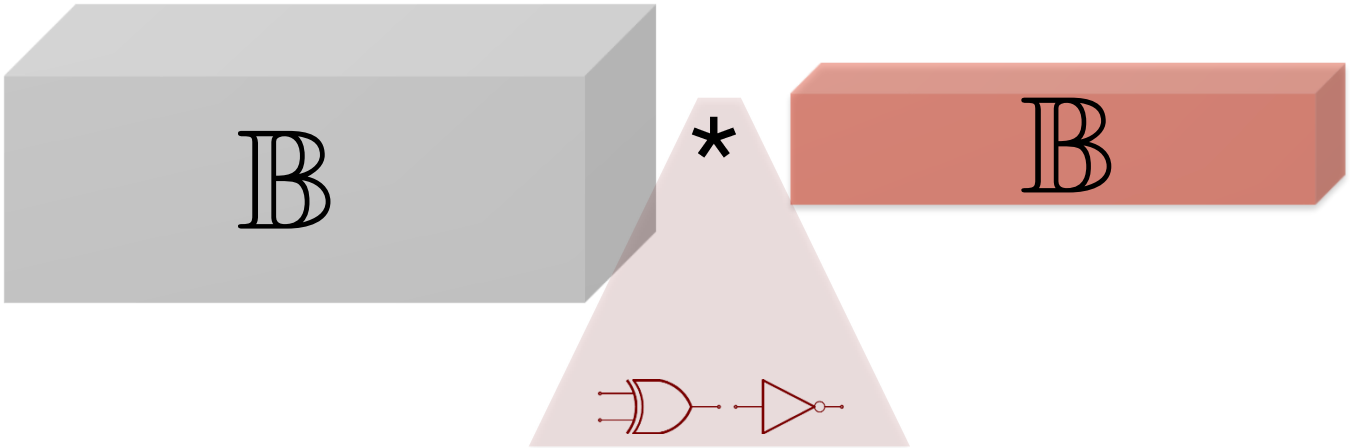

Training Binary Weight Networks

Naive Solution:

1. Train a network with real value parameters
2. Binarize the weight filters

XNOR.AI Technology

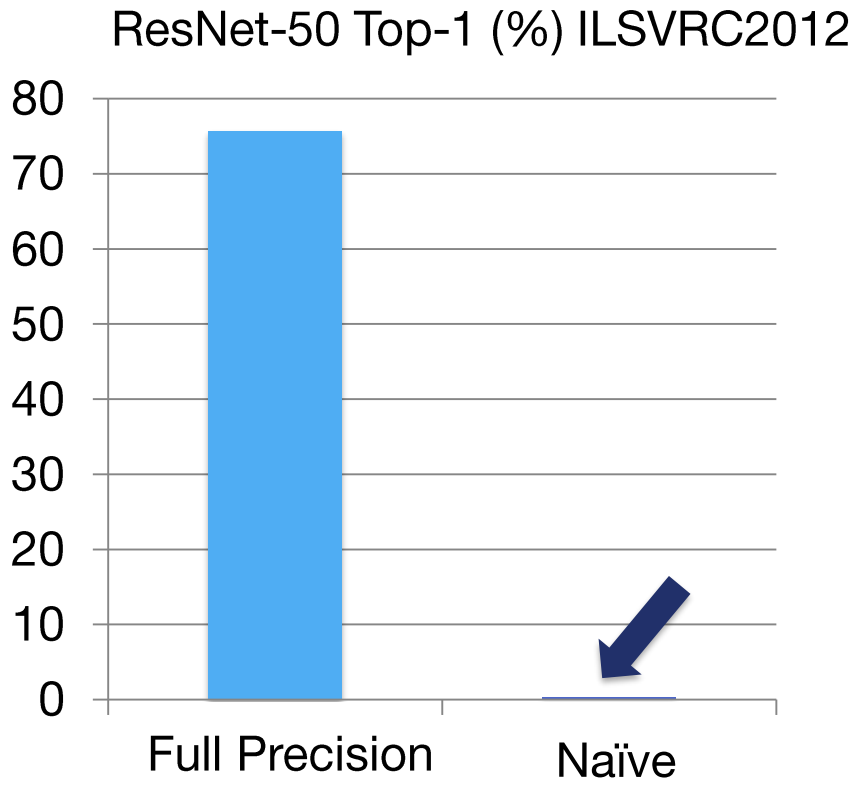

XNOR.AI Technology

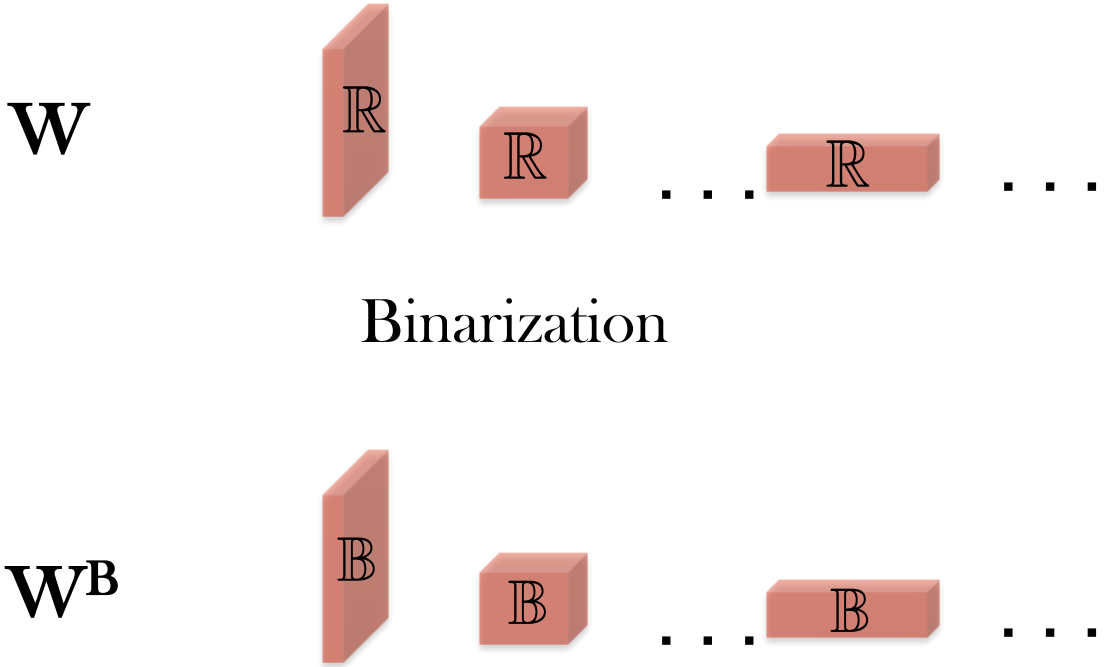

XNOR.AI Technology

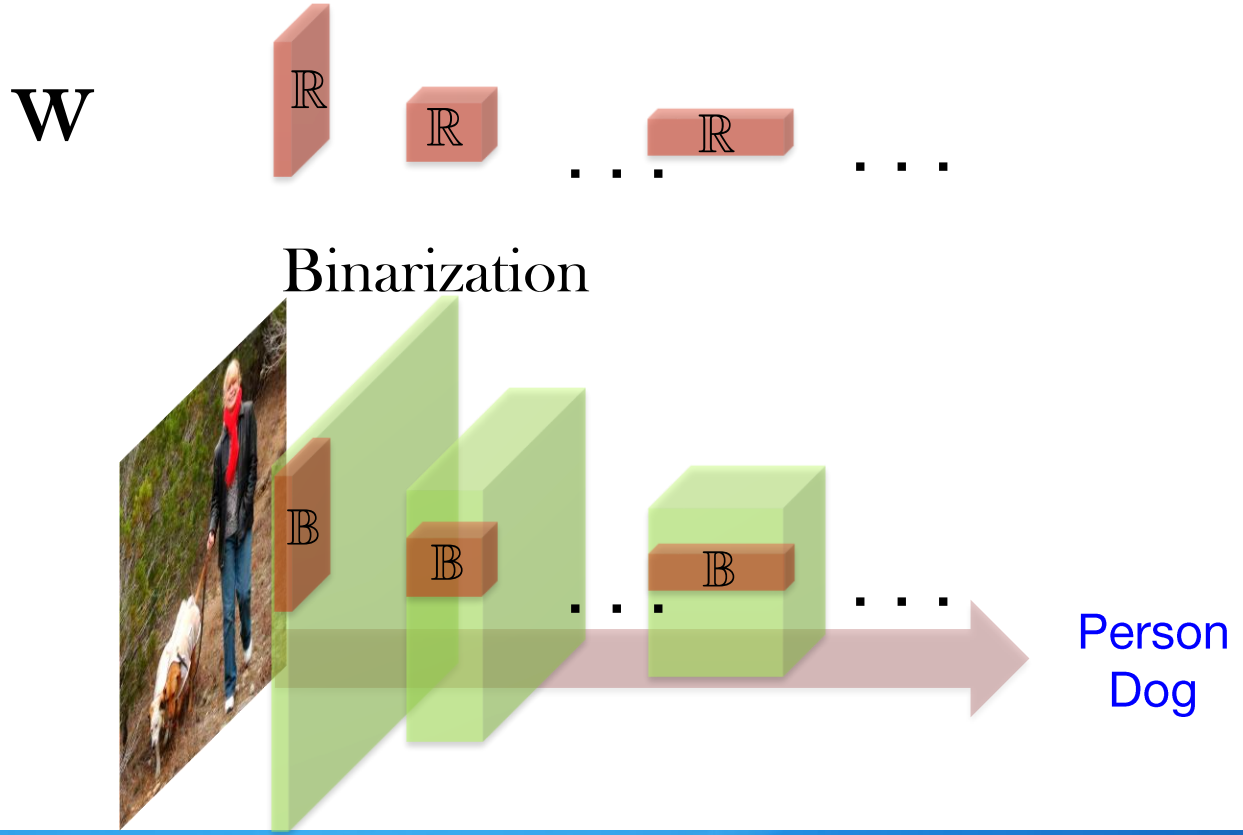

XNOR.AI Technology

Train for binary weights:

1. Randomly initialize $\mathbf{W}$
2. For $iter = 1$ to N
3. Load a random input image $\mathbf{X}$
4. $\mathbf{W^B} = \text{sign}(\mathbf{W})$
5. $\alpha = \frac{\|W\|_{\ell 1}}{n}$
6. Forward pass with $\alpha, \mathbf{W^B}$
7. Compute loss function $\mathbf{C}$
8. $\frac{\partial \mathbf{C}}{\partial \mathbf{W}}$ = Backward pass with $\alpha, \mathbf{W^B}$
9. Update $\mathbf{W}$ $(\mathbf{W} = \mathbf{W} - \frac{\partial \mathbf{C}}{\partial \mathbf{W}})$

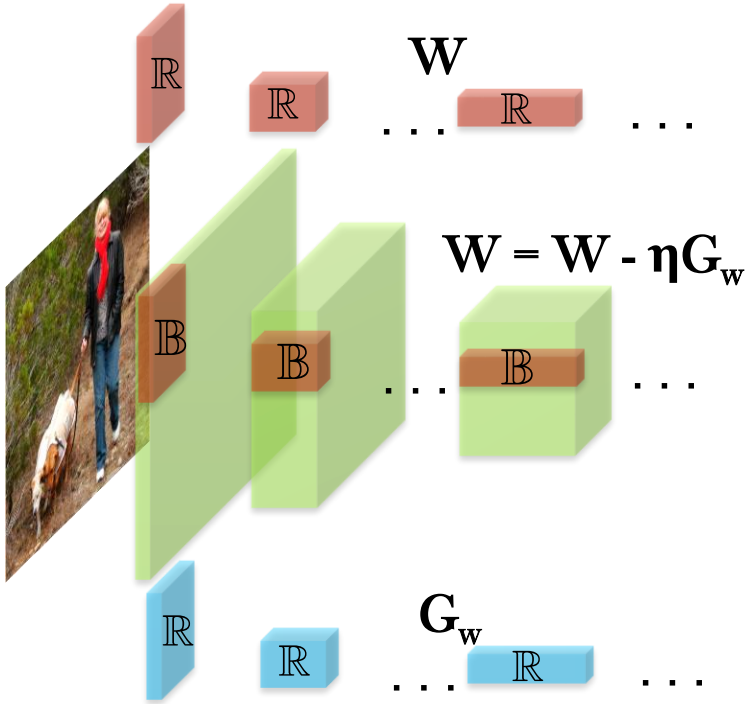

XNOR.AI Technology

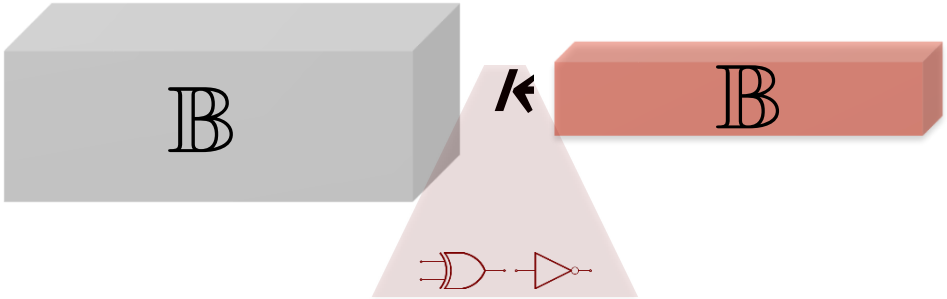

1. Randomly initialize $\mathbf{W}$
2. For $iter = 1$ to N
3. Load a random input image $\mathbf{X}$
4. $\mathbf{W^B} = \text{sign}(\mathbf{W})$
5. $\alpha = \frac{\|W\|_{\ell 1}}{n}$
6. Forward pass with $\alpha, \mathbf{W^B}$
7. Compute loss function $\mathbf{C}$
8. $\frac{\partial \mathbf{C}}{\partial \mathbf{W}}$ = Backward pass with $\alpha, \mathbf{W^B}$
9. Update $\mathbf{W}$ $(\mathbf{W} = \mathbf{W} - \frac{\partial \mathbf{C}}{\partial \mathbf{W}})$

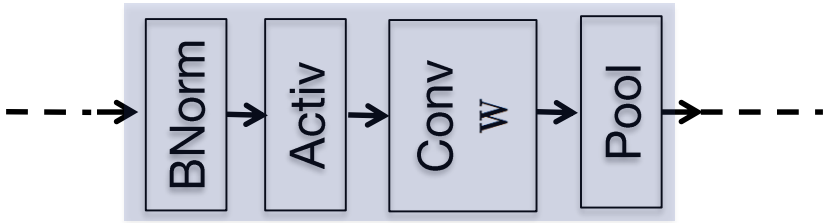

XNOR.AI Technology

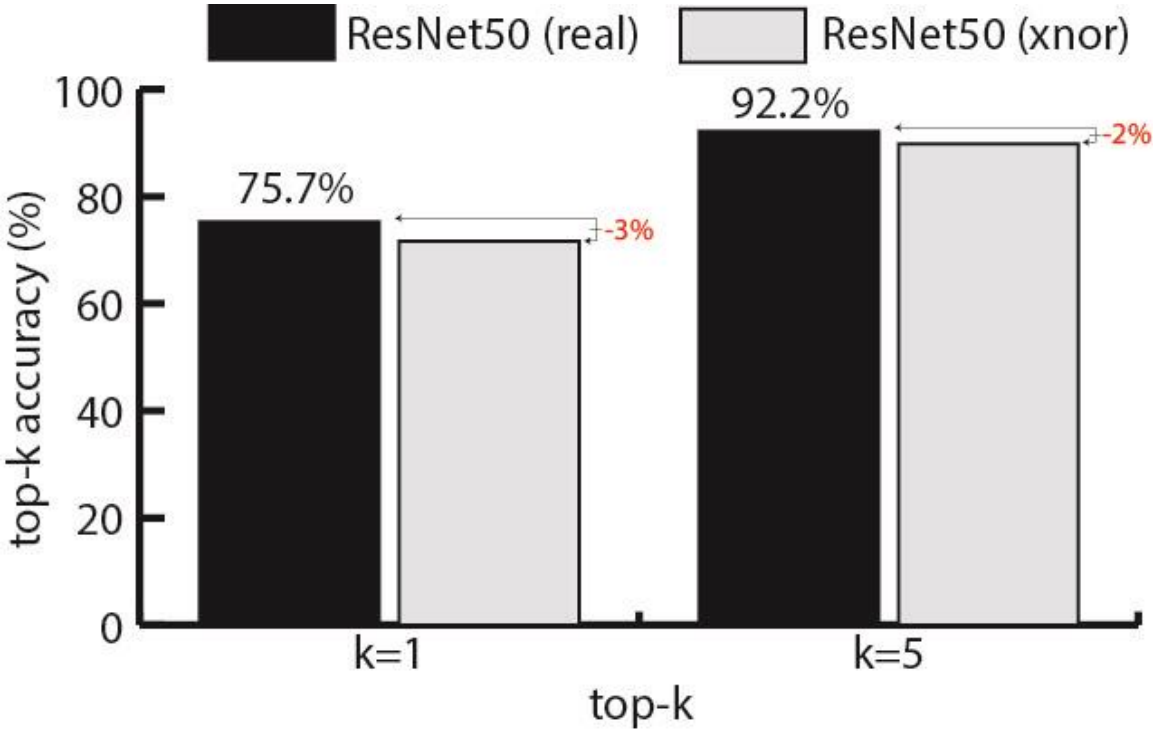

- 15x Smaller
- 10x Faster
- 200% power efficiency

XNOR.AI Technology

XNOR.AI

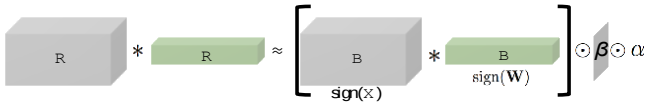

1. Randomly initialize W
2. For *iter* = 1 to *N*
3. Load a random input image X
4. W^B = sign(W)
5. $\lambda = \frac{kW_k}{n}$
6. Forward pass with λ, W^B
7. Compute loss function C
8. $\frac{\partial C}{\partial W}$ = Backward pass with λ, W^B
9. Update W ($W = W - \frac{\partial C}{\partial W}$)

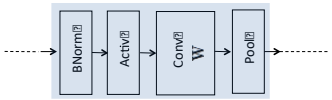

Xnor.ai IP

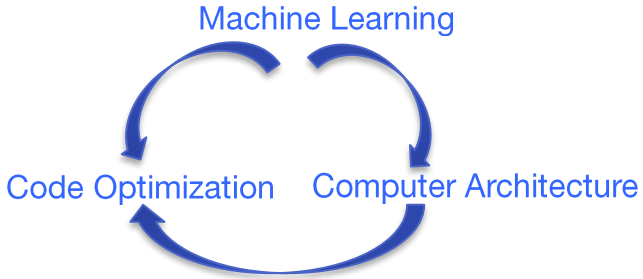

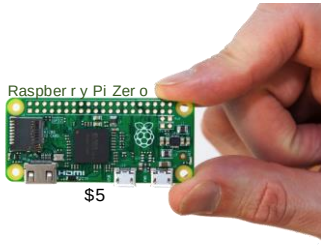

Integrate with XNOR.AI

How to integrate with XNOR.AI?

Enterprise

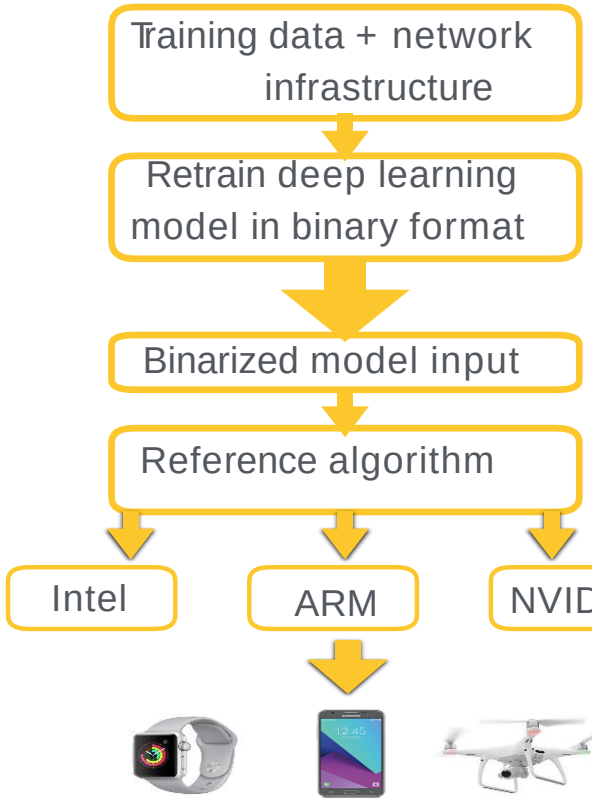

XNOR.AI deep learning Deep learning platform for model creation and optimization

XNOR.AI CORE

A proprietary code base that automatically translates binarized deep learning models into highly optimized programs that run on the leading hardware platforms

Product integrations for devices

Custom work done as part of our services agreement

Developers:

AI Everywhere

For Everyone

Integrate with XNOR.AI

Developers Platform

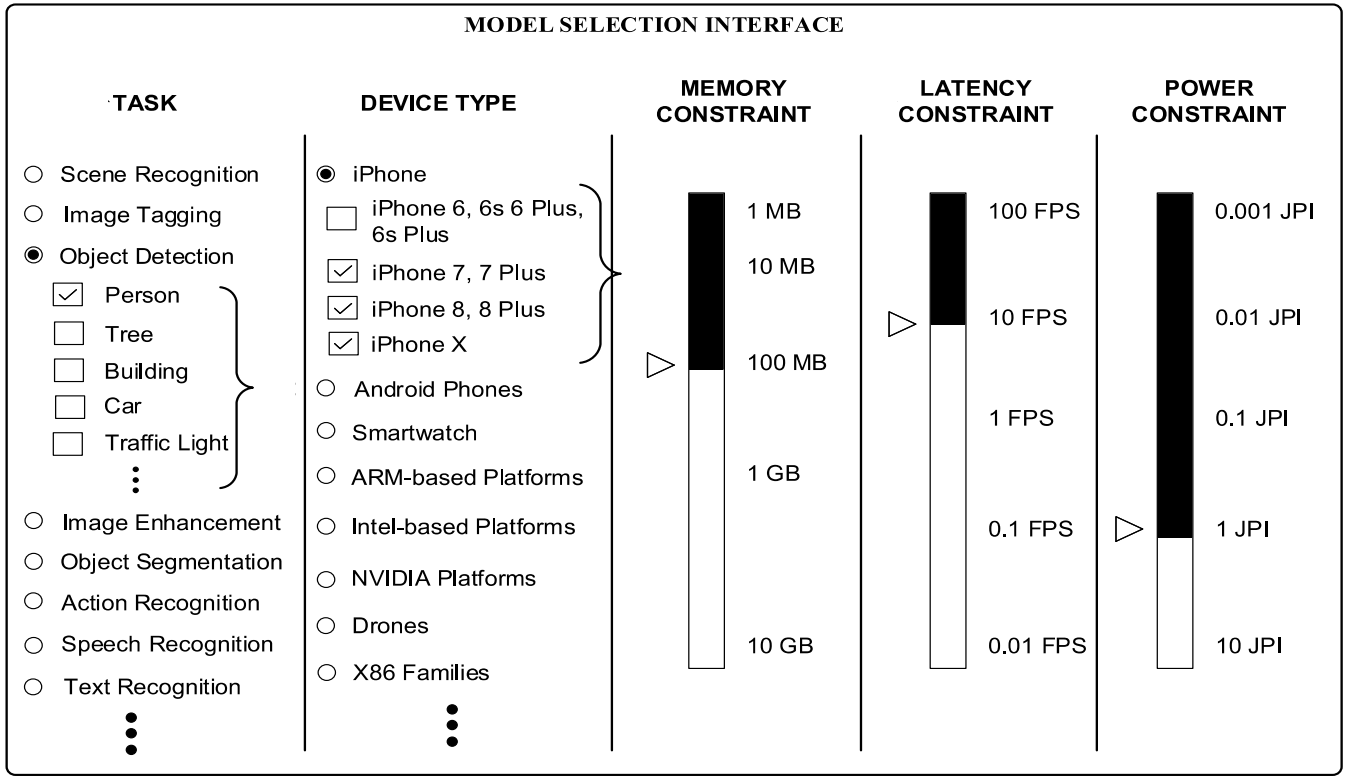

XNOR.AI technology powers multiple domains

- ✓ Aerospace & Surveillance
- ✓ Driver Assisted Systems
- ✓ Retail
- ✓ Consumer Mobile

Company Background

AI Everywhere

Founded:
2017 by Professor Ali Farhadi
and Dr. Mohammad Rastegari

Intellectual Property:
Highly strategic patent portfolio covering efficient AI at the edge

Press:
New York Times
Tech Crunch

Board members:
Ali Farhadi (CEO), Oren Etzioni (CEO of Allen Institute for AI), Matt McIlwain (Madrona Venture Group)

XNOR Innovations in AI:
XNOR-Net, Yolo, Yolo9000, LCNN, Neural Speed Reading, understanding actions (imSitu), question answering (BiDAF)

Awards:
Best paper at CVPR 2017
CVPR 2016 People's Choice Award

Thank you !!!

Learn more

Visit our table #809

www.xnor.ai

Mohammad Rastegari | Chief Technology Officer | mohammad@xnor.ai