

From Feature Engineering to Network Engineering

ShatterLine Labs

Auro Tripathy May 22, 2018

Outline



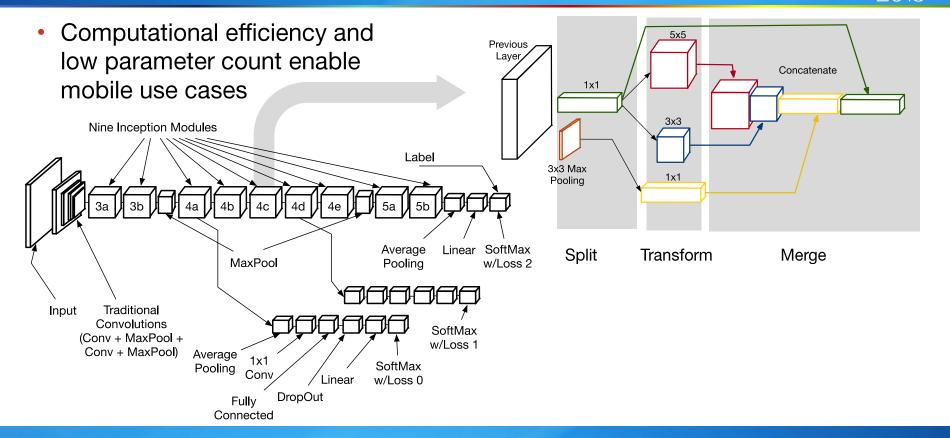
- Review Four Network Topologies
 - Split-transform-merge connections for reduced parameters
 - Residual connections for higher accuracy through depth
 - Encode/Decode networks for segmentation
 - Siamese networks for learning similarity
- Canonical Code for each Topology
- Goal Equip you with the Basic Building Blocks for your own Designs



Split-Transform-Merge Topologies to Reduce Compute



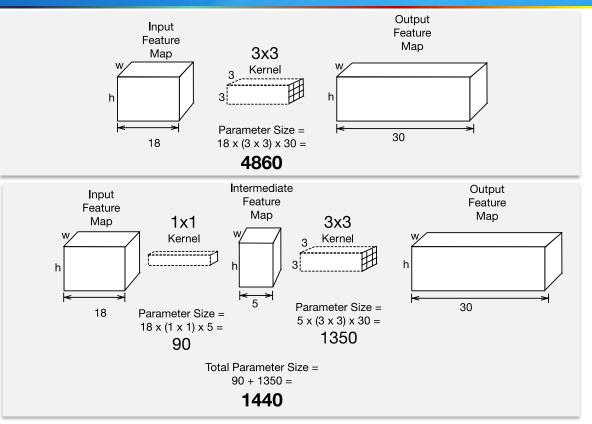
Split-Transform-Merge in the Inception Module in GoogleNet



Replacing a 3x3 Conv with a 1x1 Conv followed by a 3x3 Conv Results in Far Fewer Parameters



Same Size Input Feature Maps



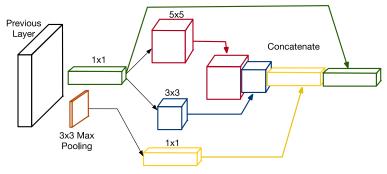
Same Size
Output
Feature
Maps

Parameters reduce from 4860 to 1440

Expressing in Keras



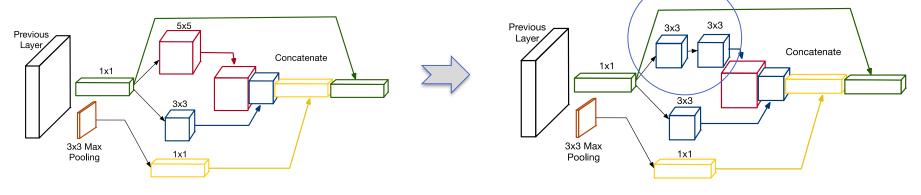
```
conv_1x1 = Conv2D(192, (1, 1), padding='same', activation='relu')(previous_layer)
```

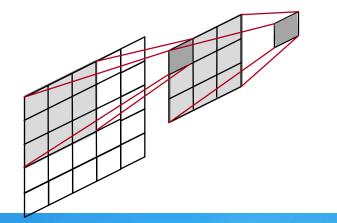


pool = MaxPooling2D(pool_size=(3,3),strides=(1,1), padding='same') (previous_layer)
conv 1x1 after pool = Conv2D(64, (1, 1), padding='same', activation='relu') (pool)

Layered Factorization into Smaller, Less Expensive Convolutions



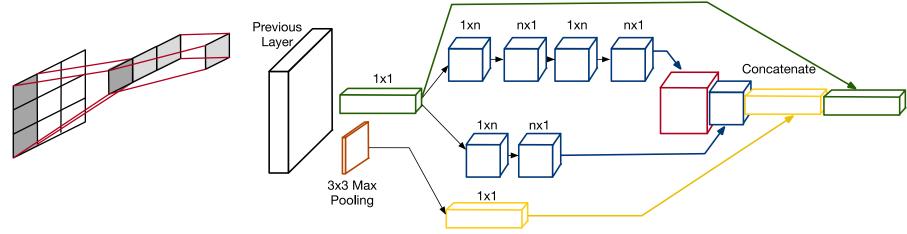




- Replace 5x5 convolution with two layers of 3x3 convolutions
- Reduces parameter count from 25 to 18
- "Note that we have factorized the traditional 7x7 convolution into three 3x3 convolutions..."

Spatial Factorization into Separable Kernels





- Factored into two convolutions a 3x1 convolution followed by a 1x3 convolution
 - Equivalent to sliding a two layer network with the same receptive field as in a 3X3 convolution
 - 33% cheaper (6 versus 9 parameters)

Resource for Inception Networks



- Paper
 - Going deeper with convolutions
 - Christian Szegedy et. Al.
- Code
 - https://github.com/keras-team/keras/tree/master/keras/applications

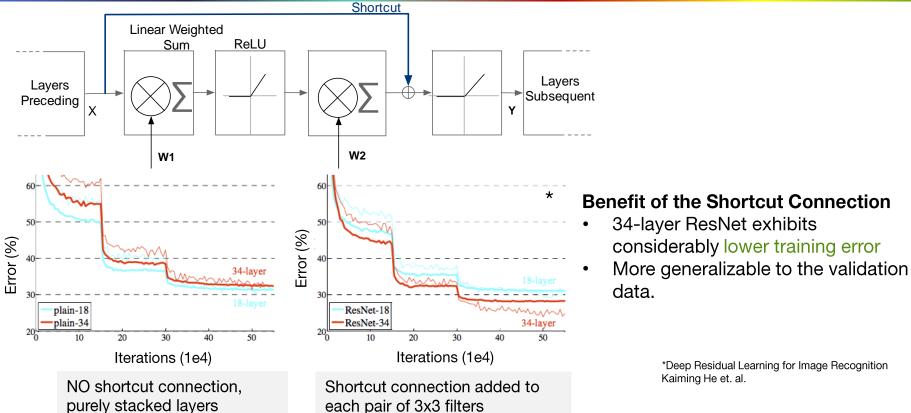


Residual Topologies for Simplicity and Uniformity



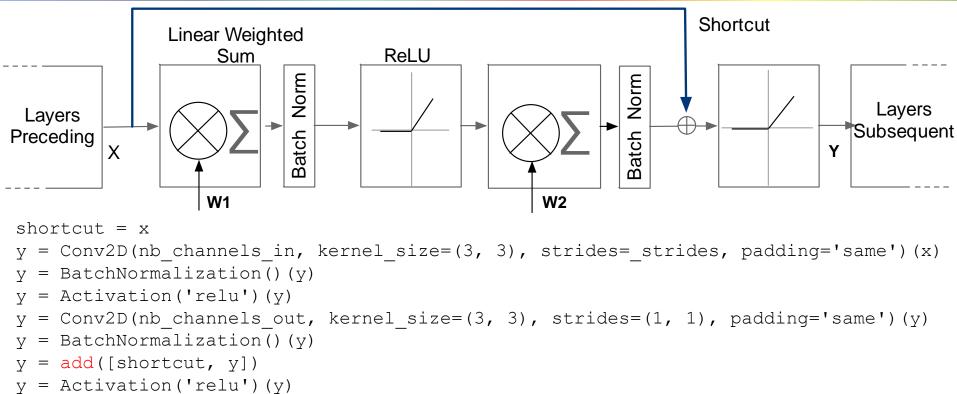
Shortcut Connections Lead to Accuracy Gains with Increased Depth





Expressing the Residual Block in Keras





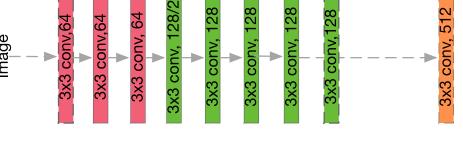
High-level View of Residual Networks (ResNets)



Average Pool

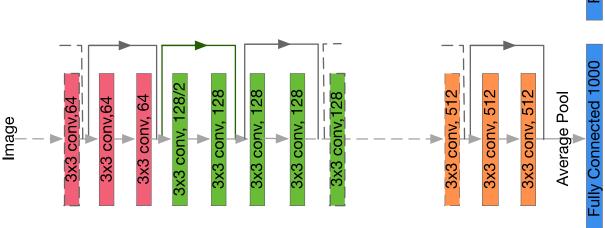
Pure Stacked Network:

Higher Error with
Increased Depth



ResNet - Stacked + Shortcut Connections:

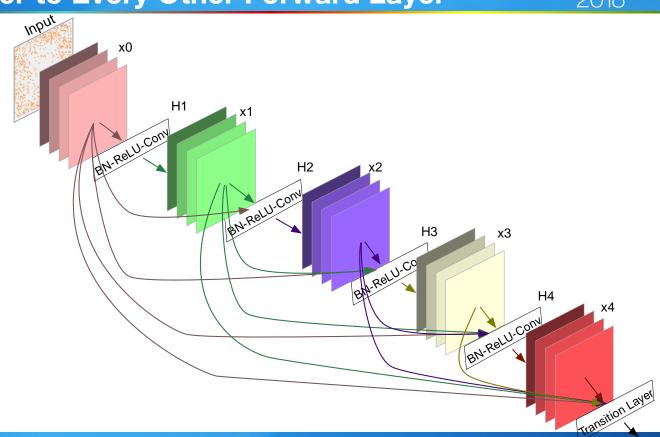
Lower Error with Increased Depth



DenseNet: Connects Each Layer to Every Other Forward Layer



- Each Layer takes all preceding featuremaps as input
- Like ResNet, however skip connections amplified
- Unlike ResNet, features are concatenated, not summed.
- Requires one-third the parameters of ResNet



Expressing a Five-Layer Dense Block with a Growth-Rate of k=4



```
def dense block(x,
                nb filters,
                growth rate,
                nb conv layers,
                dropout fraction=None,
                weight decay=1E-4):
    feature map list = [x]
    for i in range (nb conv layers):
        x = composite bn relu conv(x,
                          growth rate,
                                   dropout fraction,
                                            weight decay)
        feature map list.append(x)
        x = Concatenate(axis=-1(feature map list))
        nb filters += growth rate
    return nb filters, x
```

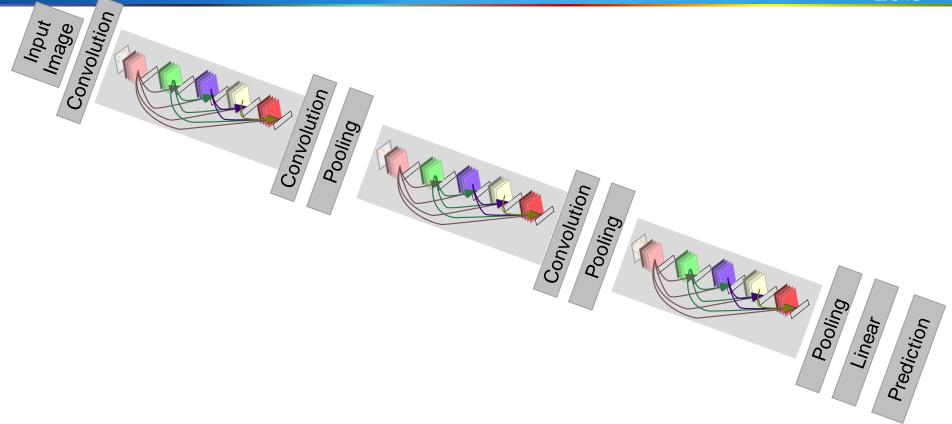
Expressing the Composite Function Batch Norm, followed by ReLU, followed by 3X3 Conv



```
def composite bn relu conv(x,
                          nb filters,
                          dropout fraction=None,
                          weight decay=1E-4):
    x = BatchNormalization(axis=1,
             beta regularizer=12 (weight decay),
             gamma regularizer=12(weight decay))(x)
    x = Activation('relu')(x)
    x = Convolution2D(nb filters,
                      kernel size=(3,3),
                      padding='same',
                      use bias=False,
                      kernel initializer='he normal',
                      kernel regularizer=12(weight decay))(x)
    if dropout fraction:
        x = Dropout(dropout fraction)(x)
    return x
```

Putting Together DenseNet





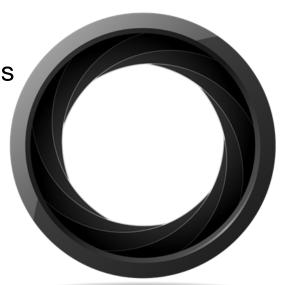
Resources for Resnet and DenseNet



- Paper
 - Deep Residual Learning for Image Recognition
 - Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun
 - Densely Connected Convolutional Networks
 - Gao Huang, Zhuang Liu, Laurens van der Maaten
- Code
 - https://github.com/kerasteam/keras/blob/master/keras/applications/resnet50.py
 - https://github.com/flyyufelix/DenseNet-Keras



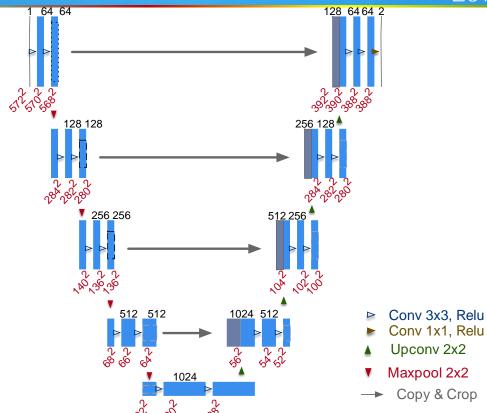
Encode/Decode Topologies for Segmenting Objects



UNET: A Contracting Path Followed by Expanding Path for Segmenting Images

embedded VISION SUMMIT 2018

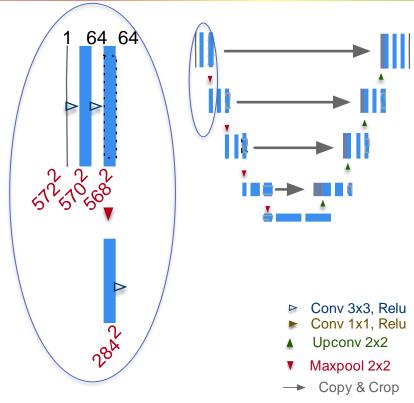
- Left side performs downsampling (a.k.a. max pooling operations) lowering resolution
- Right side performs upsampling operations increasing resolution ending up with the same resolution
- To segment, at each level the down-sampled features are combined with upsampled output followed by a convolution.



Left Side Convolutions and Pooling - One



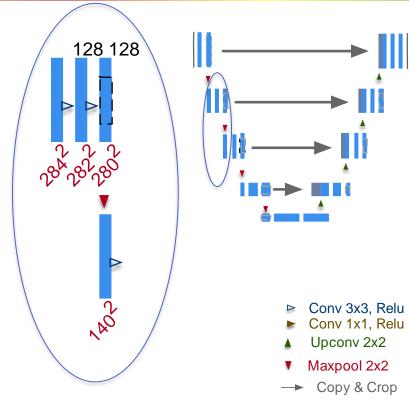
```
conv1 L = Conv2D(64, (3, 3),
       activation='relu',
       padding='same') (inputs)
conv1 L = Conv2D(64, (3, 3),
       activation='relu',
       padding='same') (conv1 L)
pool1 = MaxPooling2D(
       pool size=(2, 2)) (conv1 L)
```



Left Side Convolutions and Pooling - Two



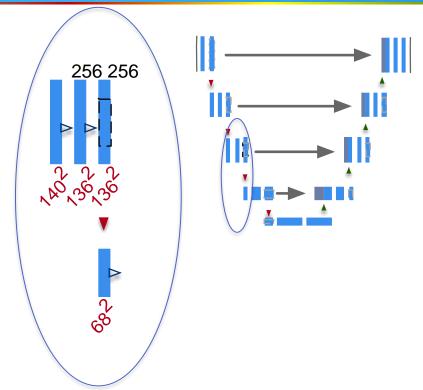
```
conv2 L = Conv2D(128, (3, 3),
       activation='relu',
       padding='same') (pool1)
conv2 L = Conv2D(128, (3, 3),
       activation='relu',
       padding='same') (conv2 L)
pool2 = MaxPooling2D(
       pool size=(2, 2)) (conv2 L)
```



Left Side Convolutions and Pooling - Three



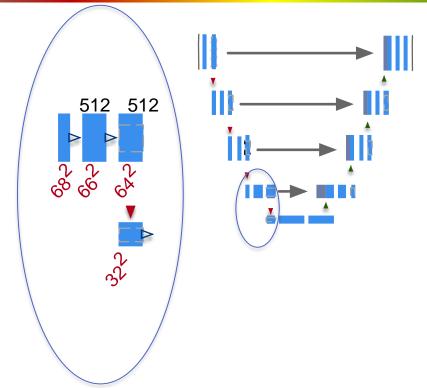
```
conv3 L = Conv2D(256, (3, 3),
       activation='relu',
       padding='same') (pool2)
conv3 L = Conv2D(256, (3, 3),
       activation='relu',
       padding='same') (conv3 L)
pool3 = MaxPooling2D(
       pool size=(2, 2)) (conv3 L)
```



Left Side Convolutions and Pooling - Four

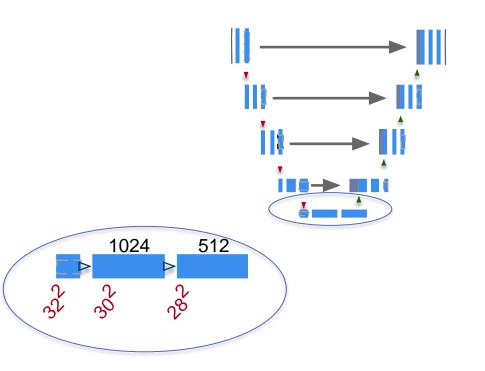


```
conv4 L = Conv2D(512, (3, 3),
       activation='relu',
       padding='same') (pool3)
conv4 L = Conv2D(512, (3, 3),
       activation='relu',
       padding='same') (conv4 L)
pool4 = MaxPooling2D(
       pool size=(2, 2)) (conv4 L)
```



Bottom Convolutions





Concatenation, Up Sampling, and Convolutions - One



```
up1 = concatenate(
        [Conv2DTranspose (512, (2, 2),
       strides=(2, 2),
       padding='same') (conv5 B),
conv4 L],
       axis=3)
conv3 R = Conv2D(256, (3, 3),
                                                 1024 512
       activation='relu',
       padding='same') (up1)
conv3 R = Conv2D(256, (3, 3),
       activation='relu',
                                            2 512
       padding='same') (conv3 R)
```

Concatenation, Up Sampling, and Convolutions - Two



```
up2 = concatenate(
        [Conv2DTranspose (512, (2, 2),
        strides=(2, 2),
       padding='same') (conv3 L),
conv3 L],
                                                512 256
        axis=3)
conv2 R = Conv2D(256, (3, 3),
        activation='relu',
       padding='same') (up2)
conv2 R = Conv2D(256, (3, 3),
        activation='relu',
        padding='same') (conv2 R)
```

Concatenation, Up Sampling, and Convolutions - Three

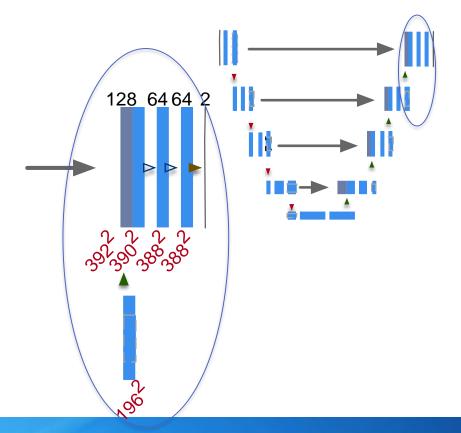


```
up3 = concatenate(
        [Conv2DTranspose (128, (2, 2),
       strides=(2, 2),
       padding='same') (conv2 R),
                                              256 128
conv2 L],
       axis=3)
conv1 R = Conv2D(128, (3, 3),
       activation='relu',
       padding='same') (up3)
conv1 R = Conv2D(128, (3, 3),
       activation='relu',
       padding='same') (conv1 R)
```

Concatenation, Up Sampling, and Convolutions - Four

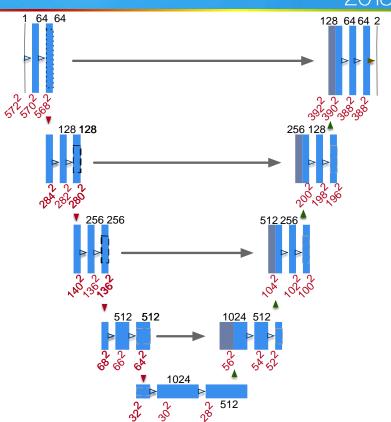


```
up4 = concatenate(
        [Conv2DTranspose (64, (2, 2),
       strides=(2, 2),
       padding='same') (conv1 R),
conv1 L],
       axis=3)
conv out = Conv2D(64, (3, 3),
       activation='relu',
       padding='same') (up3)
conv out = Conv2D(64, (3, 3),
       activation='relu',
       padding='same') (conv out)
```



Final Segmentation





Resources for Unet



- Paper
 - U-Net: Convolutional Networks for Biomedical Image Segmentation
 - Olaf Ronneberger, Philipp Fischer, and Thomas Brox
- Code
 - https://github.com/zhixuhao/unet



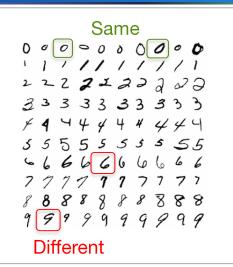
Siamese CNN for Learning Similarities



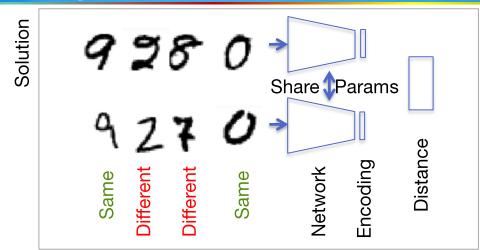
Problem Statement Decide Whether Pairs of Images are Similar/Different



Problem



- 1. Generate image pairs
 - 1. For same digits
 - 2. For different digits



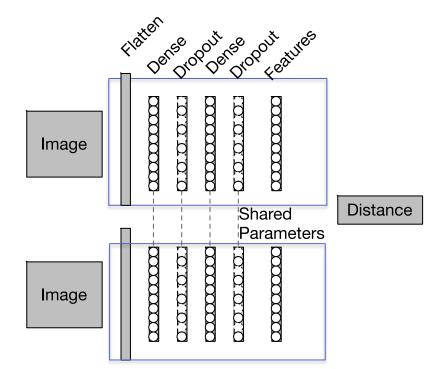
- Learn parameters such that :
 - 2. For same digits, the distance is small
 - 3. For different digits, the difference is large
- 3. Can be done by applying gradient descent on
 - 2. Triplet Loss, or
 - 3. Contrastive Loss
- 4. Uses-cases include (1) Signature verification (2) Face Verification

Siamese Base Network – Identical Branches



Code for each branch

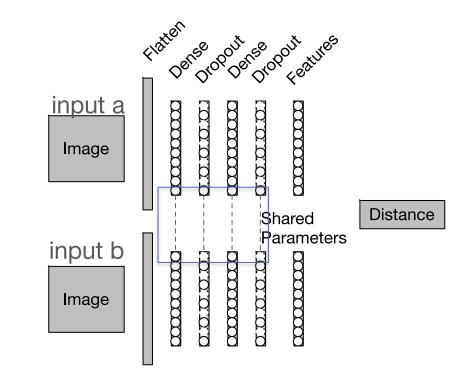
```
input = Input(shape=input_shape)
x = Flatten()(input)
x = Dense(128, activation='relu')(x)
x = Dropout(0.1)(x)
x = Dense(128, activation='relu')(x)
x = Dropout(0.1)(x)
x = Dense(128, activation='relu')(x)
```



Share Parameters Across the Two Branches



```
siamese a = base network(input a)
siamese b = base network(input b)
distance =
       Lambda (euclidean distance,
                   output shape=
       eucl dist output shape)
        ([siamese a, siamese b])
model = Model([input a, input b],
distance)
```



Resources for Siamese Networks



Papers

- DeepFace: Closing the Gap to Human-Level Performance in Face Verification
 - · Yaniv Taigman, Ming Yang, Marc'Aurelio Ranzato, Lior Wolf
- FaceNet: A Unified Embedding for Face Recognition and Clustering
 - Florian Schroff, Dmitry Kalenichenko, James Philbin

Code

 https://github.com/kerasteam/keras/blob/master/examples/mnist_siamese.py

Summary



- The shift from better features to better network topologies
- We reviewed four:
 - Split-transform-merge connections for reduced parameters
 - Residual connections for higher accuracy thru depth
 - Encode/Decode networks for segmentation
 - Siamese networks for learning similarity