

# embedded **VISION** SUMMIT 2018

## Deploying CNN based vision solutions on a \$3 microcontroller

**Au-Zone Technologies** is a leading provider of development tools and enabling IP used for the design of intelligent embedded vision products and solutions.

Our architecture agnostic development tools (eCV SDK and DeepView) enable our customers quickly to develop and securely deploy machine learning solutions on a range of embedded hardware.

## Design Challenge

---




# Toaster Oven Example

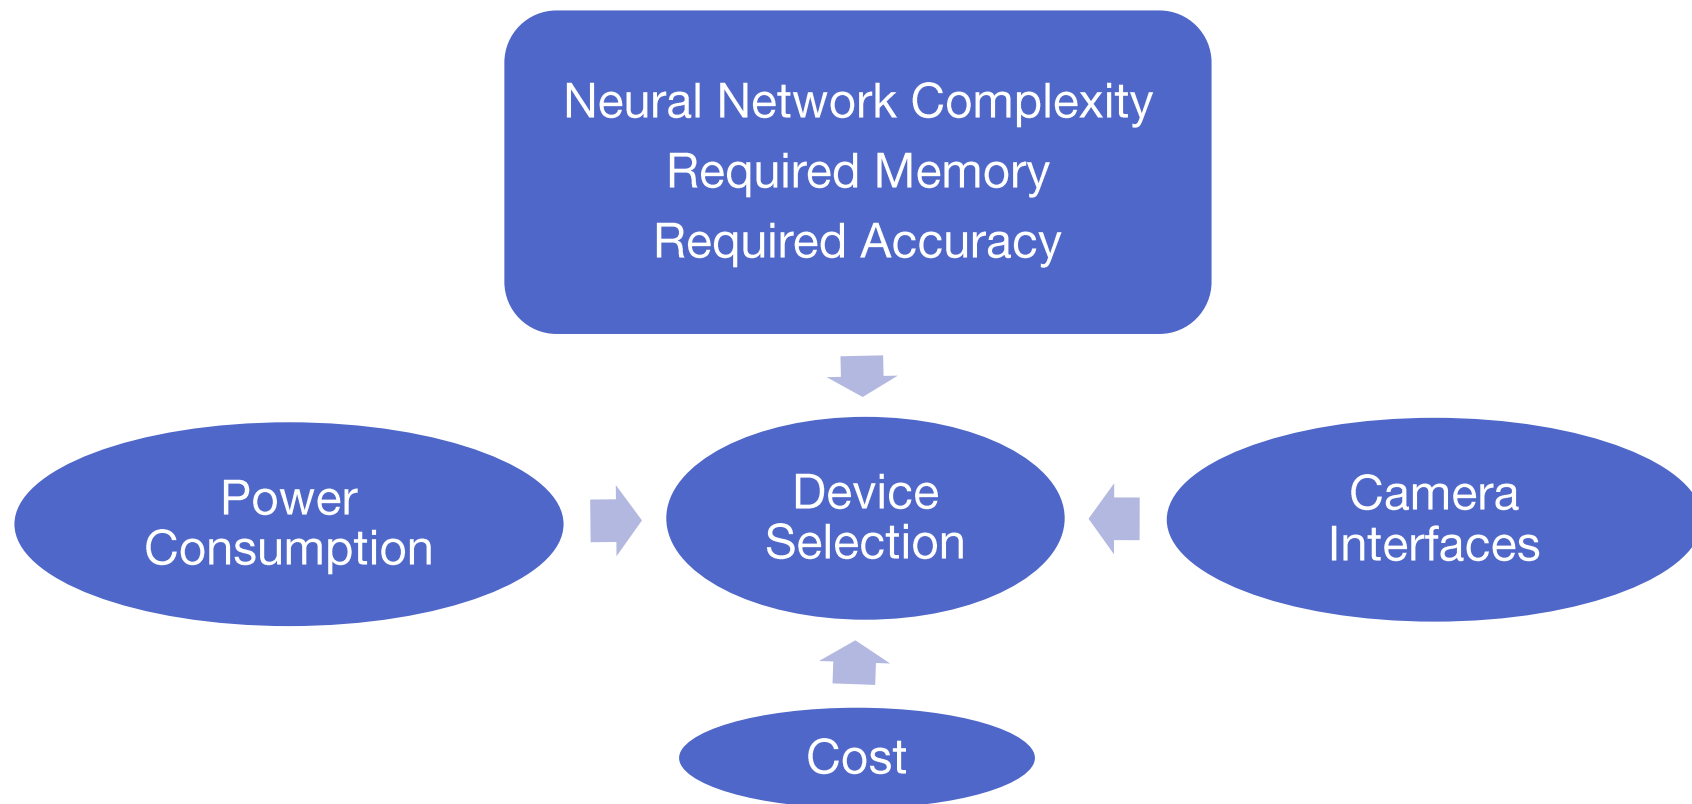


## Requirements:







- Embedded image capture and classification
- Image Sensor
- LCD & Buttons
- Compute & Memory for vision algorithm and application
- Low BOM cost



# Design Considerations



# Architecture Options





Many Options?

# Market Drivers

- Is this practical?
- Why should you care?

Image Classification enabled




Low Incremental Cost

# Constraints of Cortex M vs Cortex A

Much Lower raw  
processor performance

Significant Memory  
Constraints

## CPU Performance & Memory






# Network Design


---




# Network Model Design Flow



# Network Layer Types




Fully Connected Layer



Convolutional Layer

# Fully Connected Operation




One Weight required for every connection

8-bit quantization saves on FLASH size with low impact on error

Parameter	Resource	Size	8bit	32 bit float	Typical
Input Buffer	RAM	6 x 6 x 32	1.1 KB	4.6 KB	4.6 KB
Weights	FLASH	6 x 6 x 32 x 1 x 1 x 75	86.4 KB	345.6 KB	86.4 KB
Output	RAM	1 x 1 x 75	.075 KB	3.00 KB	3 KB
Total Memory			87.6 KB	351 KB	7.6 KB RAM 86.4 KB FLASH
MAC Operations	CPU	(6 x 6 x 32) x (1 x 1 x 75)	86.4 K		86.4K

# Convolution Operation



3x3 Convolution x 16 channels


Parameter	Resource	Size	8bit	32 bit float	Typical
Input Buffer	RAM	26 x 26 x 3. (24x24x3 plus padding)	2.1 KB	8.1 KB	8.1 KB
3x3 Kernel	FLASH	3 x 3 x 3 x 16	0.42 KB	1.73 KB	0.42 KB
Output	RAM	24 x 24 x 16	9.2 KB	36.8 KB	36.8 KB
Total Mem			11.6 KB	46.7 KB	0.42 KB FLASH 44.9 KB RAM
MAC Ops	CPU	(24 x 24 x 3) x (3 x 3 x 16)	248.8 K		

## GTSRB –

- German Traffic Signs in the wild
- 39K Training Samples
- 12.6K Testing Samples
- 43 classes of signs
- Cropped images 24 x 24 x 3



# CNN with Fully Connected Layers



Model FLASH

**2 MB**

MAC Operations

~2M

Minimum RAM

74 KB

Weights

488K

Minimum Cache

2 KB

Layers

9

Optimum Cache


74 KB

Accuracy

94.9

(example for 32-bit float implementation)

# Reduced complexity




Model FLASH	2 M → 195 KB	MAC Operations	~2M
Minimum RAM	74 KB	Weights	488K → 48K
Minimum Cache	2 KB	Layers	9 → 8
Optimum Cache	74 KB	Accuracy	94.9 → 92.5%

(example for 32-bit float implementation)




# Fully Convolutional Network (FCN)



Model Size FLASH	109 KB	MAC Operations	~2 M
Minimum RAM	72 KB	Weights	26 K
Minimum Cache	2 KB	Layers	9
Optimum Cache	74 KB	Accuracy	93.3%

(example for 32-bit float implementation)


# Convolutional Network



Model FLASH	109 → 54 KB	MAC Operations	~2M
Minimum RAM	72 KB	Weights	26K → 13K
Minimum Cache	2 K → .57 KB	Layers	9 → 7
Optimum Cache	74 K → 37 KB	Accuracy	93.3 → 92.7%

(example for 32-bit float implementation)

# Convolutional Network FCN\_20



Model FLASH	109 → 20 KB
Minimum RAM	72 → 28 KB
Minimum Cache	2 → 0.2 KB
Optimum Cache	74 → 8 KB

MAC Operations	~2M → 222K
Weights	26 → 4K
Layers	9
Accuracy	93.3 → 91.5%

(example for 32-bit float Implementation)

## Target Implementation

---





How to test and deploy on target hardware?  
Requires optimized inference implementation

Vendor framework with generic network support

- Design flow to benchmark and iterate quickly with different network designs
- Profile tools for fitting designs
- Support for network update


Custom C code for selected network

- High level of optimization is possible
- Additional effort to prototype and implement

# Challenges to address


- Fixed memory allocation at compile time
- Available contiguous memory on target microcontroller?
- Buffer management for best RAM utilization
- Efficient image sensor pipeline
- Validation of accuracy on the target

# Fully Convolutional (FCN) – RAM Optimization



Minimum RAM is determined by largest layer for linear network  
Non-linear networks require RAM buffer management

# Image Sensor Pipeline

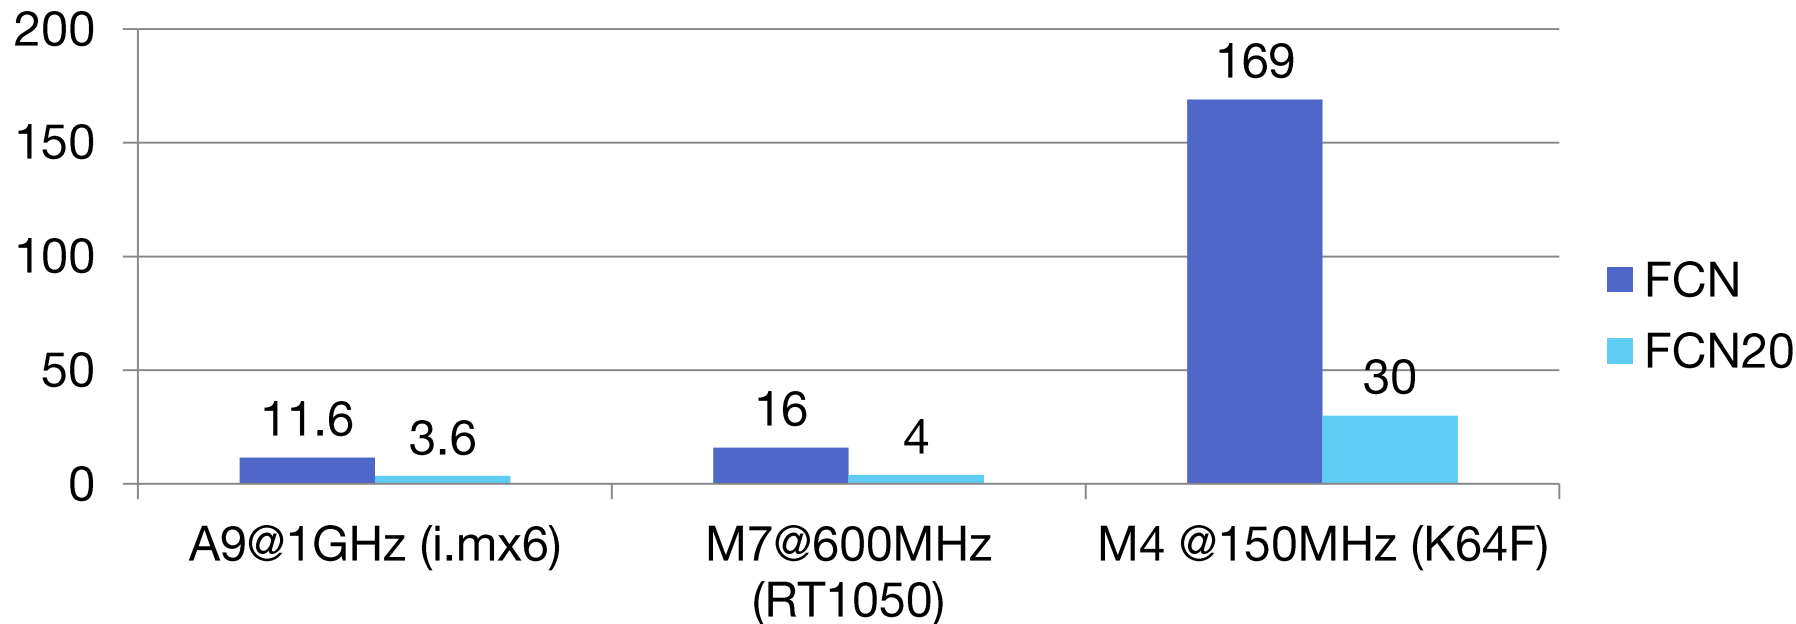


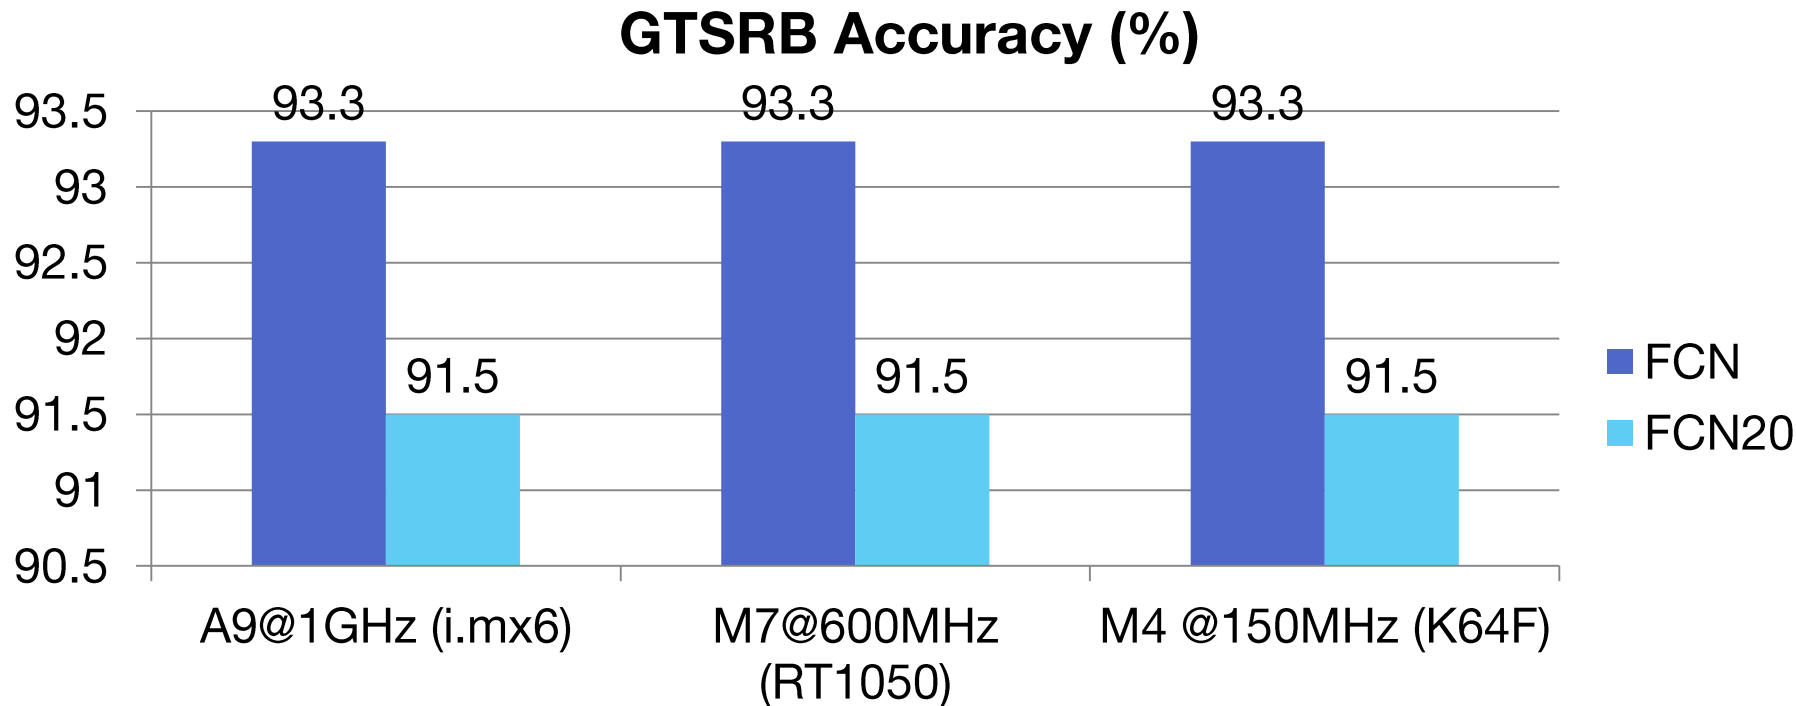
Sensor	H	V	Pixels	16bit (RGB565)
VGA	640	480	307 K	614 KB
QVGA	320	240	77 K	154 KB
QQVGA	160	120	19K	38 KB

- Use Binning or Scaling to reduce Image sensor buffer
- Use DMA to minimize slow memory buffer copies
- Use a monochrome sensor if color is not required for application



## GTSRB Inference Time (ms)






# Areas for further Optimization

- Implement networks with 8-bit /16-bit arithmetic for devices that do not have hardware FPU support
- Cloud integration for distributed edge and network-based solutions where connectivity is practical

# Validation on target demo



1. IoT-embedded vision using CNNs is practical on very low-cost microcontrollers
2. An appropriate network design is required for design and target fit
3. Memory optimization techniques should be carefully implemented
4. Specific part selection is driven by performance/accuracy and other system requirements
5. Optimized embedded inference code and tools can accelerate implementation

# Further Resources

- Drop by Au-Zone booth 802 for further demos and details
- [Software Frameworks and Toolsets for Deep Learning-based Vision Processing](#)
- [Demonstration of the DeepView ML Toolkit for Embedded Platforms](#)
- [www.embeddedml.com/deepview](http://www.embeddedml.com/deepview)
- [ARM CMSIS NN](#)
- [NXP i.MX RT Series: Crossover Processor](#)
- [NXP Kinetis® K Series:M4 Core MCU](#)

## Backup Slides

---



# Data Set Curator and Augmentation

Deep View Toolkit - Version 1.7.4-35 - C:/Users/greg.AU-ZONE/Documents/DeepView/GTSRB/GTSRB.deepview

File View Window Help

Network Model: GTSRB\_FCN

Dataset Curator Graph Designer Network Trainer Network Validator

Dataset

GTSRB

Training Samples 39209

Testing Samples 12630

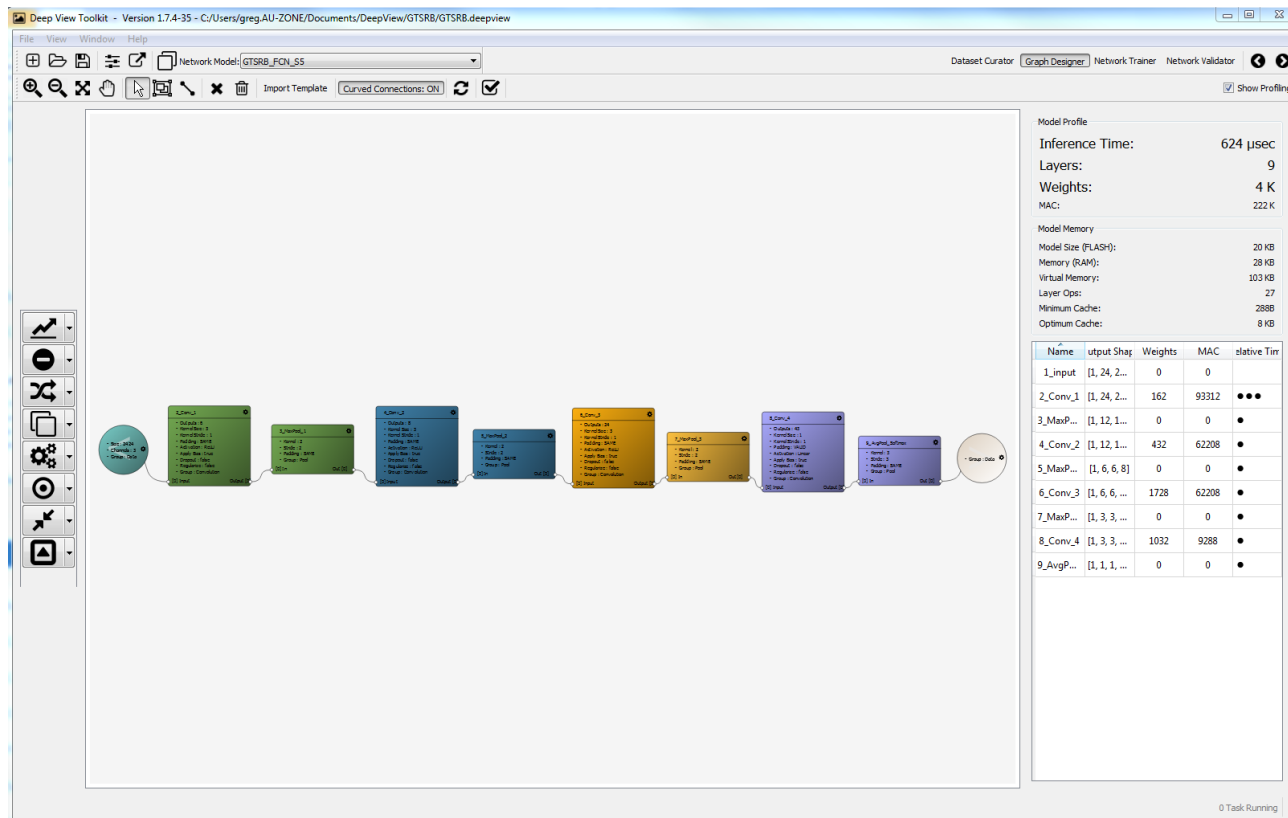
Classifications 43

Classification	# Training	# Testing
1 BewareIceSnow	450	150
2 BicycleCrossing	270	90
3 CurveLeft	210	60
4 CurveRight	360	90
5 DoNotEnter	1110	360
6 DoubleCurve	330	90
7 EndOfAllRestricti...	240	60
8 EndOfMaxSpeed80	420	150
9 EndOfNoPassing...	240	60
10 EndOfNoPassing...	240	90
11 GeneralDanger	1200	390
12 MaxSpeed100	1440	450
13 MaxSpeed120	1410	450
14 MaxSpeed20	210	60
15 MaxSpeed30	2220	720
16 MaxSpeed50	2250	750
17 MaxSpeed60	1410	450
18 MaxSpeed70	1980	660
19 MaxSpeed80	1860	630
20 MustContinueSt...	1200	300

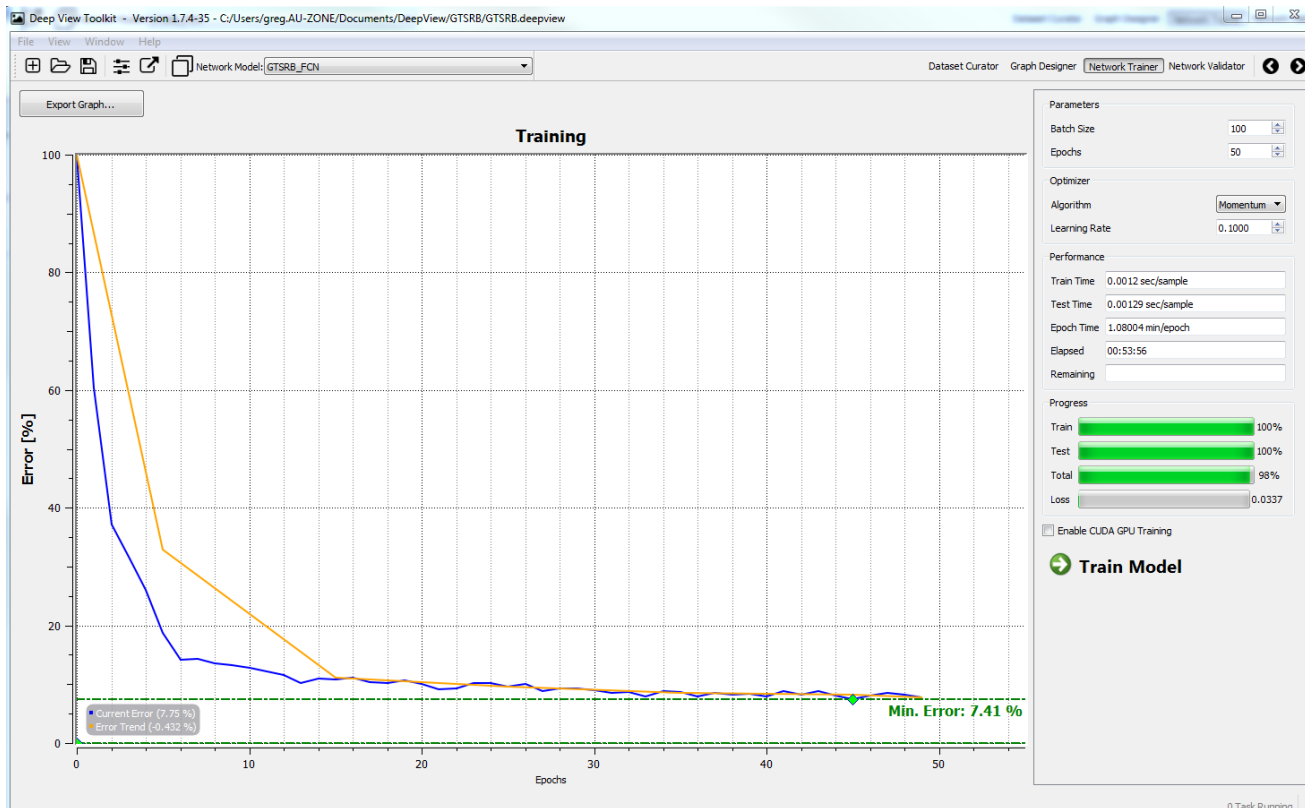
0 Task Running



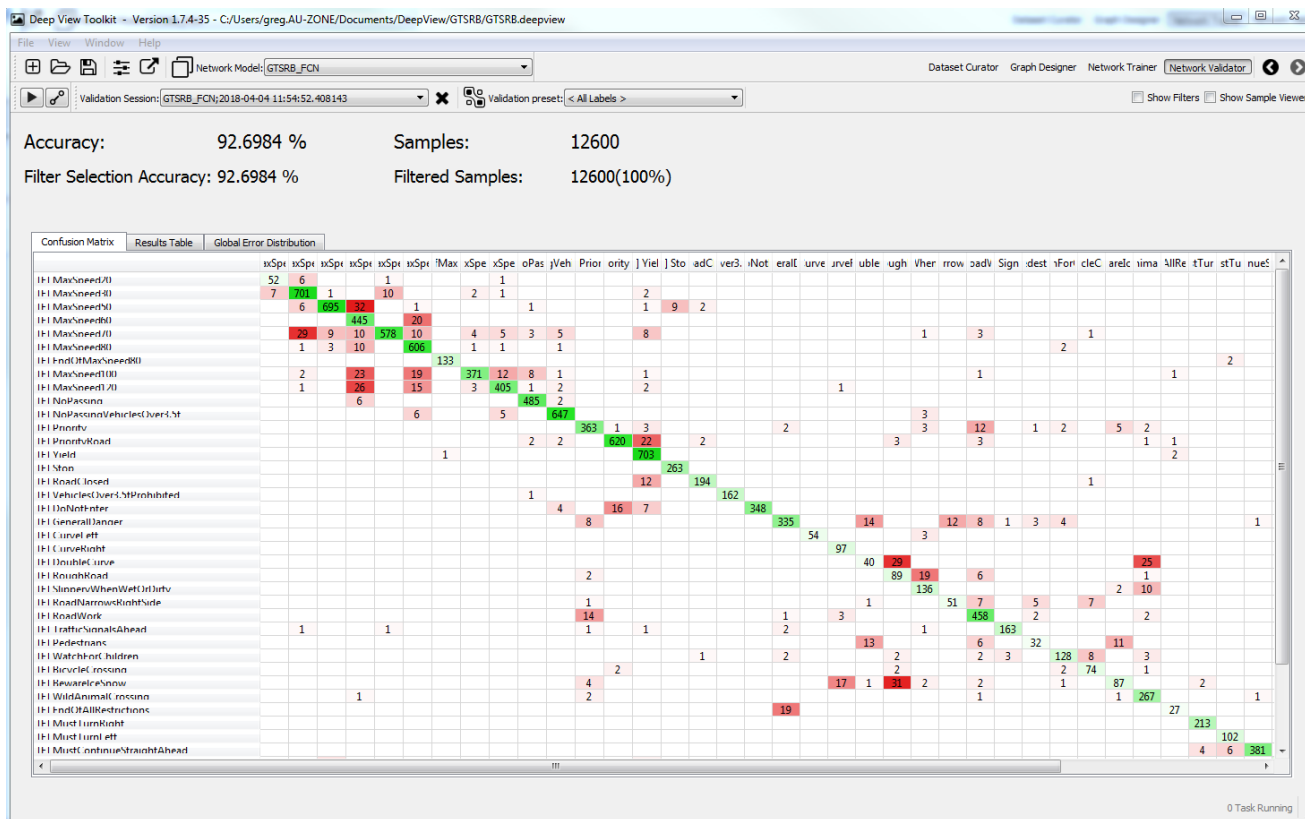
# Network Design with profile details



# Network Training



## Network Validation on Host



# Network Validation on Target

Run validation on target

Device Selection

Remote Devices FRDM64 [http://192.168.1.72:80] Add Device Remove


Validation Feedback

☒ Complete Testing Set ☐ Test Samples 1 Generate PDF Report STOP

Validation Preview Layers Inference Time

File: 251\_00251.ppm.png  
Label: Stop  
Inferred Label: Stop  
Inference Time (ms): 0

Device FRDM64 [http://192.168.1.72:80]  
Wrong Classifications 202  
Good Classifications 1294  
Accuracy 86.4973  
Inference Time (ms): 28  
Layer Ops: 27



▶

# Network Profile on Target

