Harnessing the Edge and the Cloud Together for Visual AI

Au-Zone Technologies Inc.

Sébastien Taylor

May 23, 2018

About Au-Zone

- Au-Zone Technologies is a leading provider of development tools, engineering design services, and enabling IP used for the design of intelligent embedded vision products and solutions.
- By utilizing our **Machine Learning** and **embedded Computer Vision** tools we enable our customers to quickly develop and securely deploy machine learning solutions and novel Convolutional Neural Networks on embedded hardware.

Introduction

- Focus on Image Classification using Deep Neural Networks
- Building a Hybrid Solution
- Problems to solve to make this work
  - Modeling the unknown
  - Distributing models to the edge efficiently
- Example
  - Face Recognition

Advanced Model Design

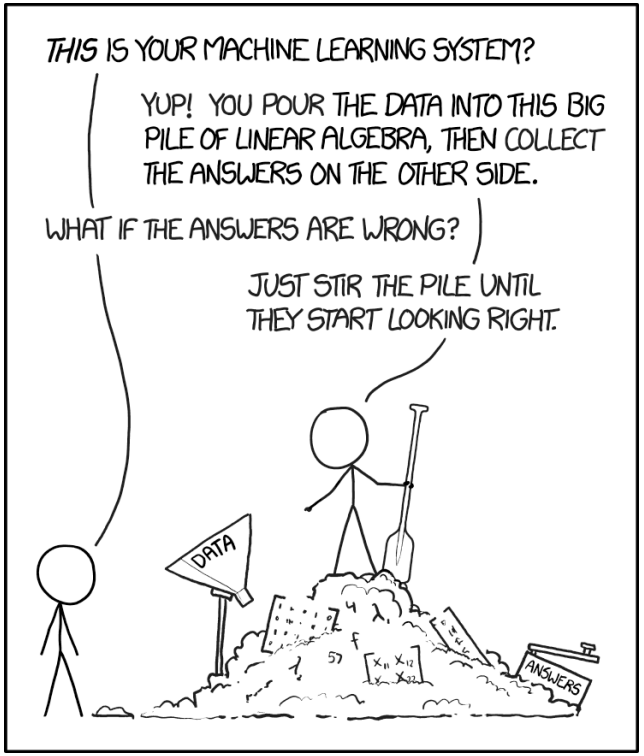

www.xkcd.org

Advanced Model Balancing

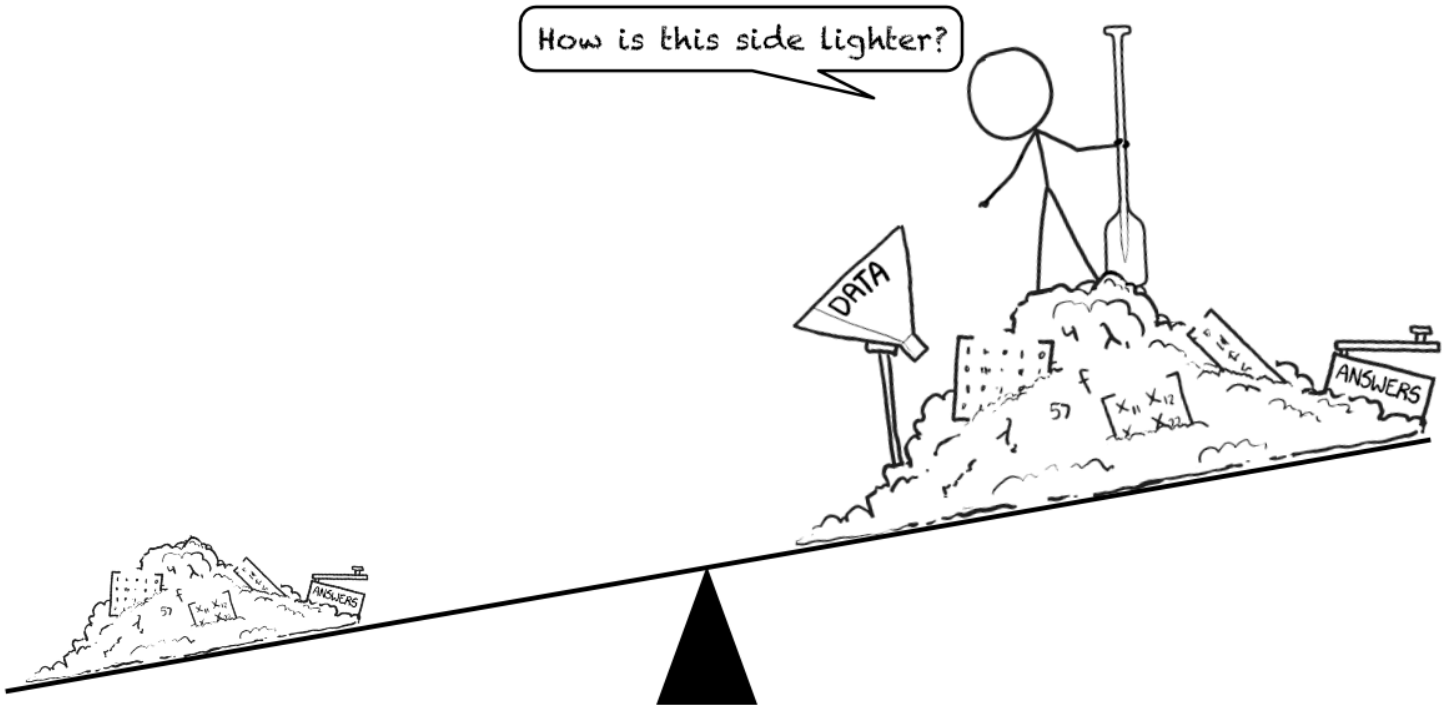

Architecture

Hybrid Edge-Cloud Architecture

- Typical Cloud Server
  - Multiple, large models
  - Central point
- Optional Edge Server
  - Intermediate between cloud and edge
  - Caching, computational offloading
  - Can handle training, dataset evolution
- Peer Nodes
  - Idle or more powerful
- Edge Nodes
  - Small models
  - Solution focused

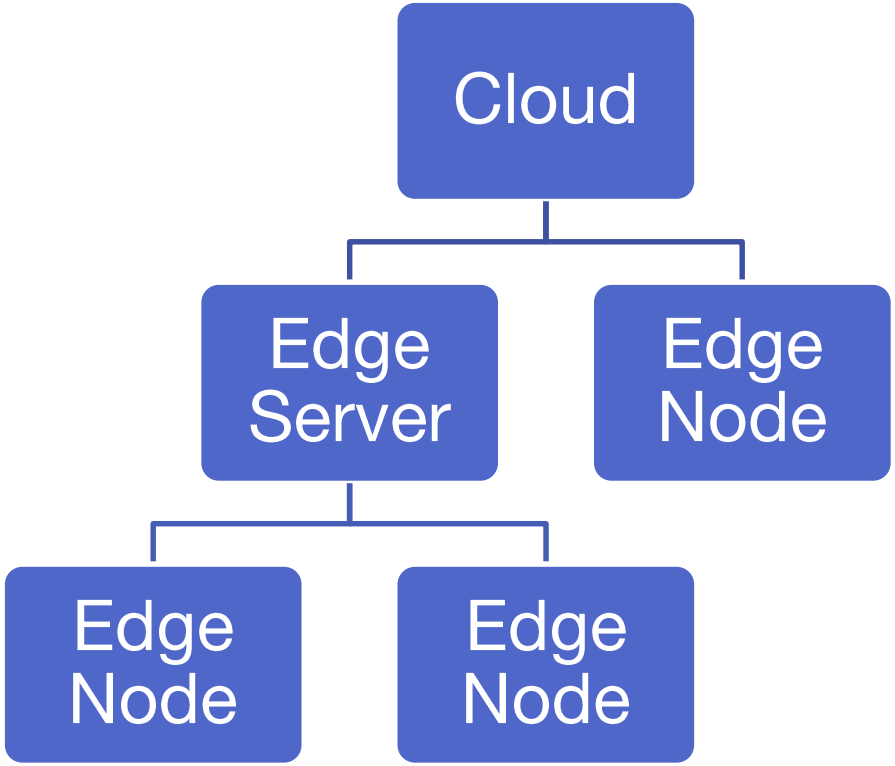

Target Edge Devices

- Our examples cover devices as small as Cortex-M4 and Cortex-M7
  - Sub-150 mW devices (CPU under $3)
  - Bare metal/RTOS
  - Hundreds of KB of RAM
- Scaling up to Cortex-A and beyond
  - Examples on Cortex-A9 and Cortex-A53 (CPU under $30)
  - Linux
  - Sub-2500 mW devices
  - Hundreds of MB of RAM

Architecture Flow

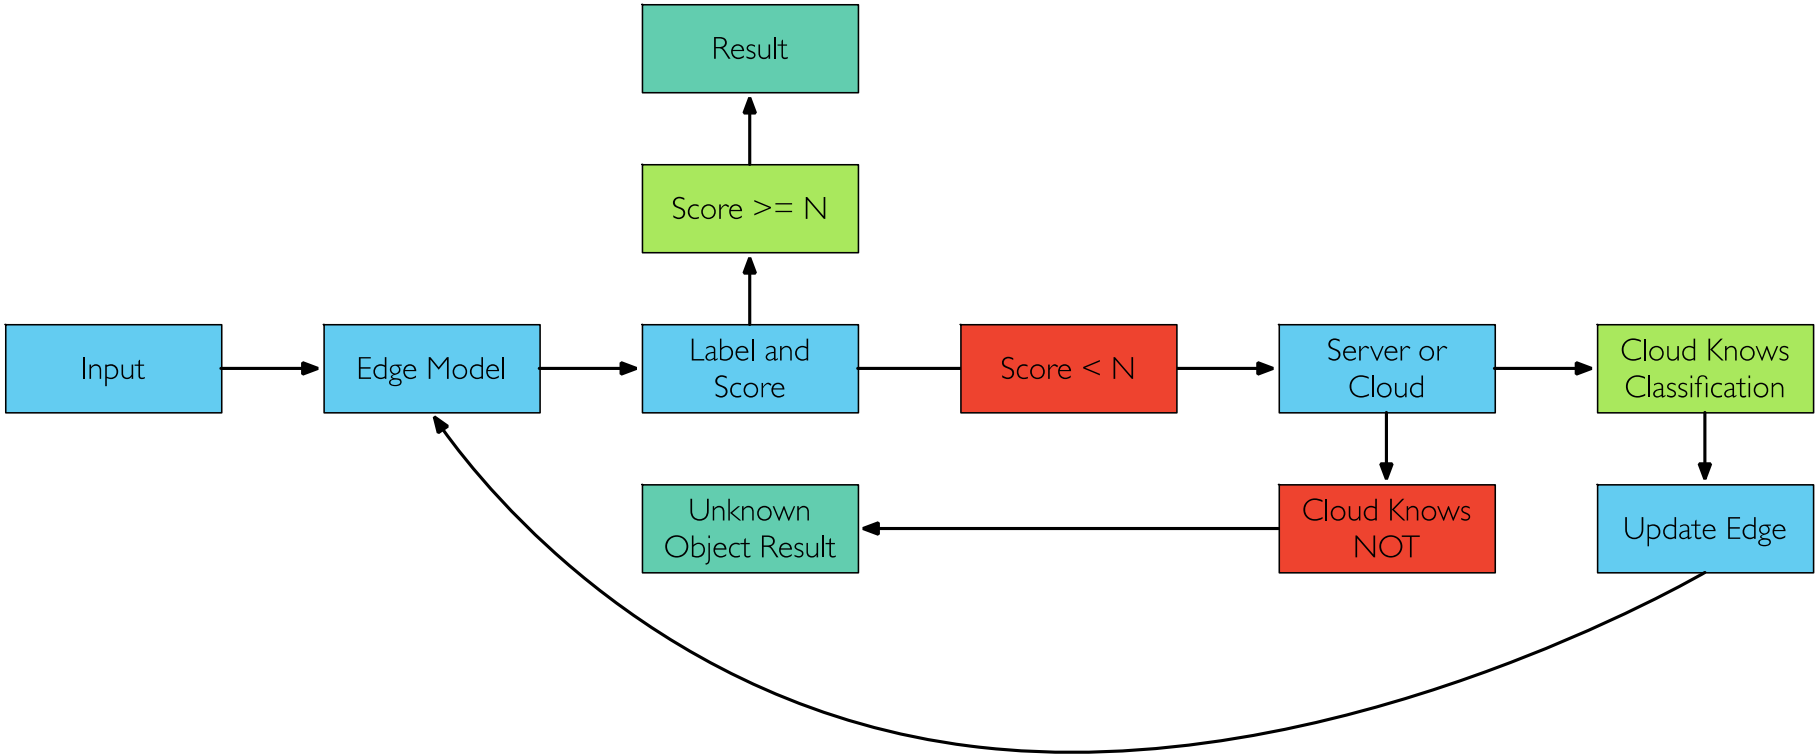

Architecture Flow – Result Scoring

How to accurately score our results?

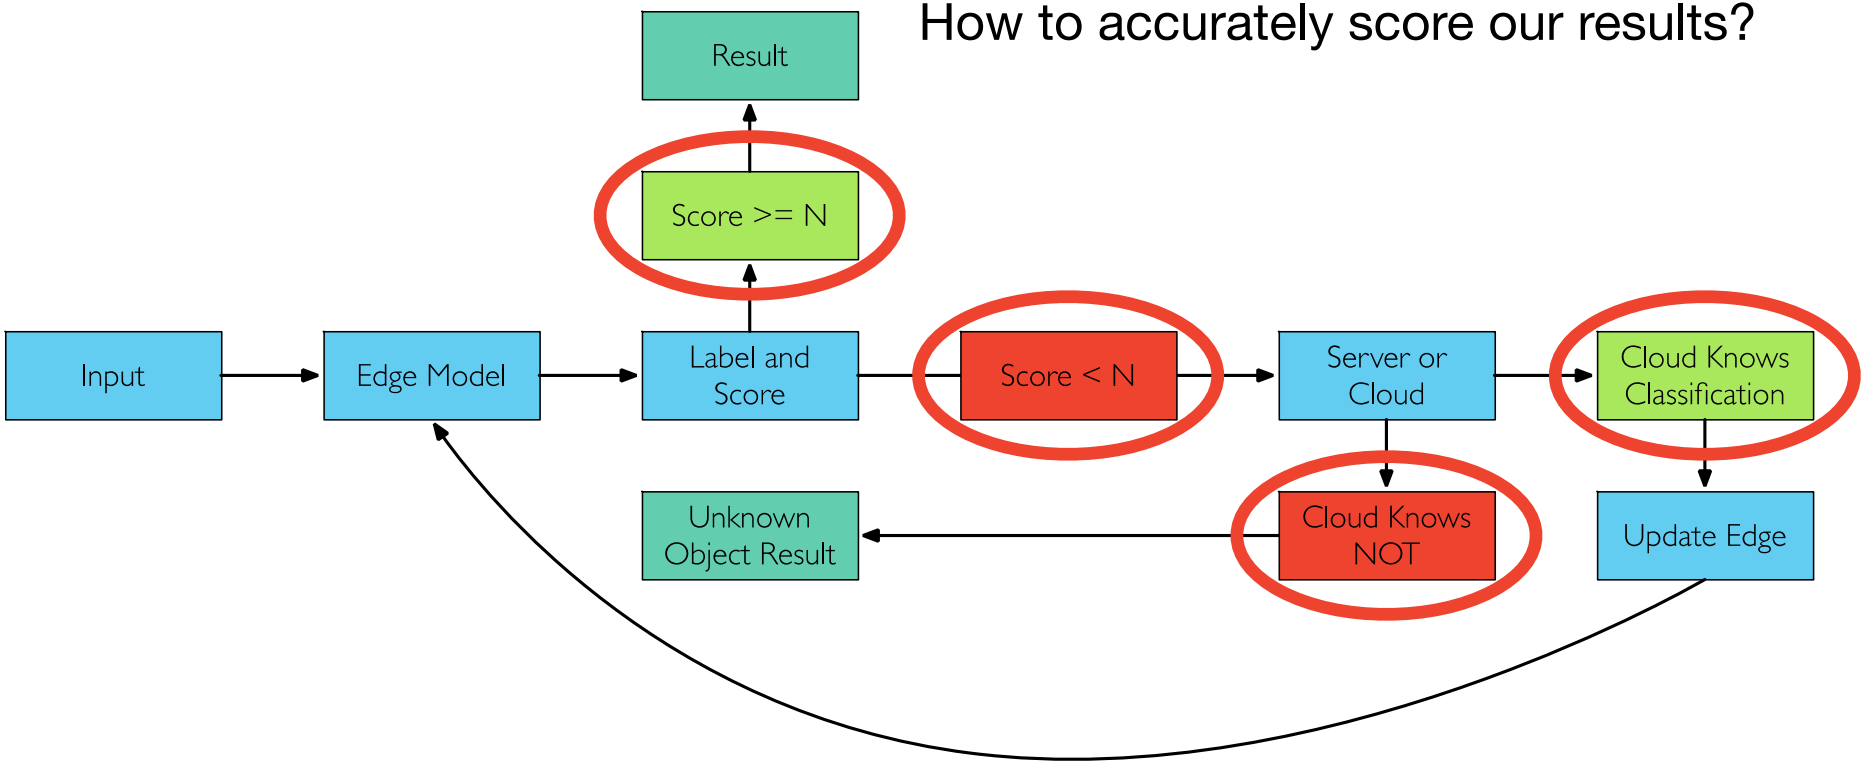

Architecture Flow – Model Update Size

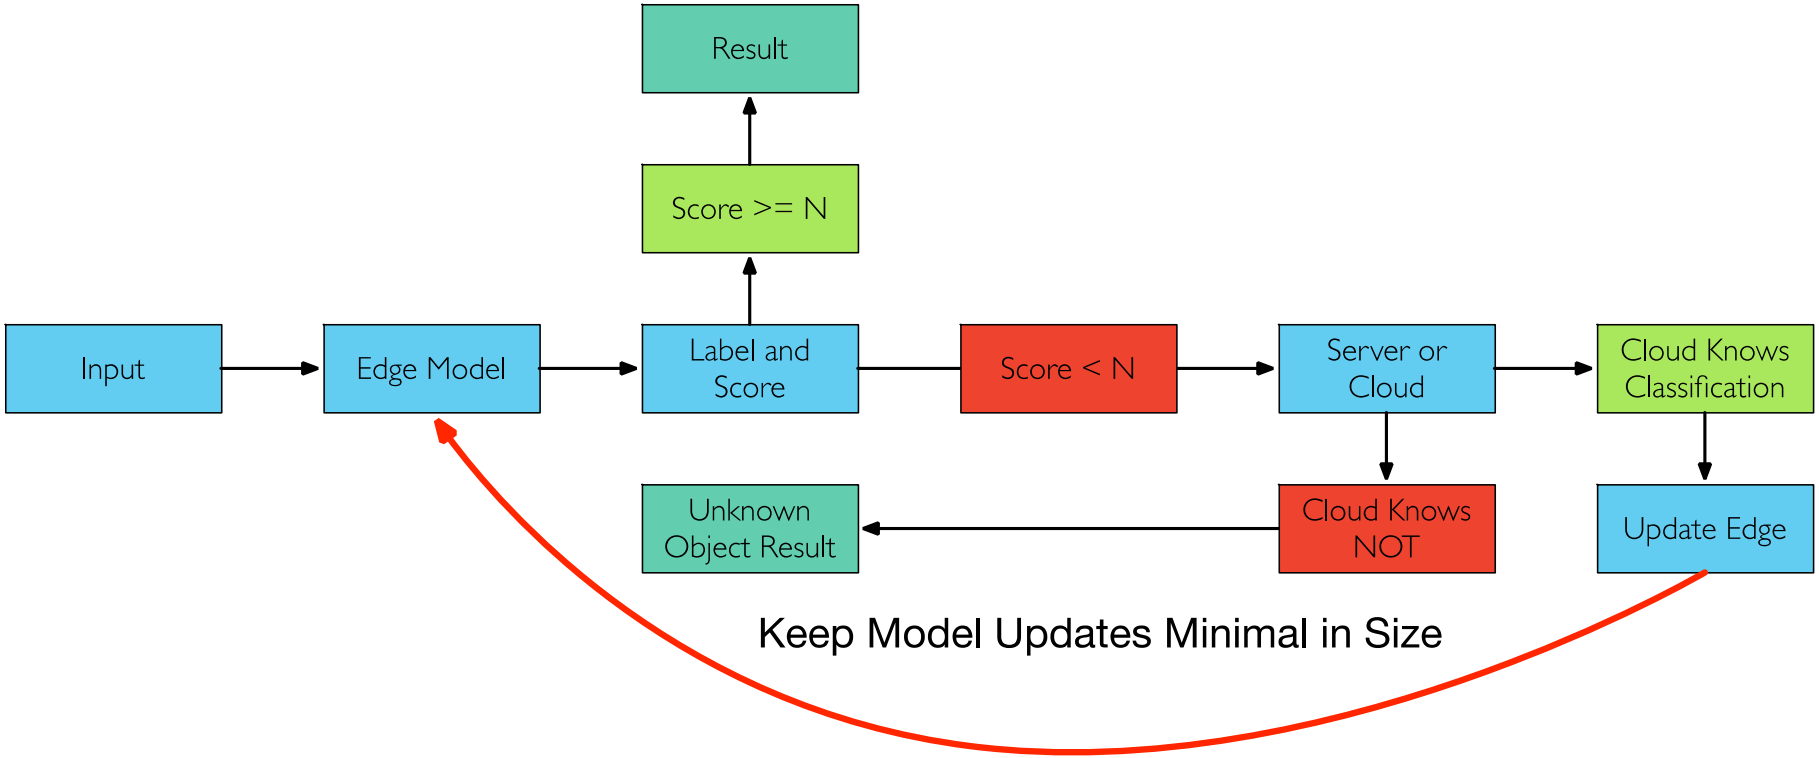

How to know what you do not know?

Architecture Flow – Result Scoring

How to accurately score our results?

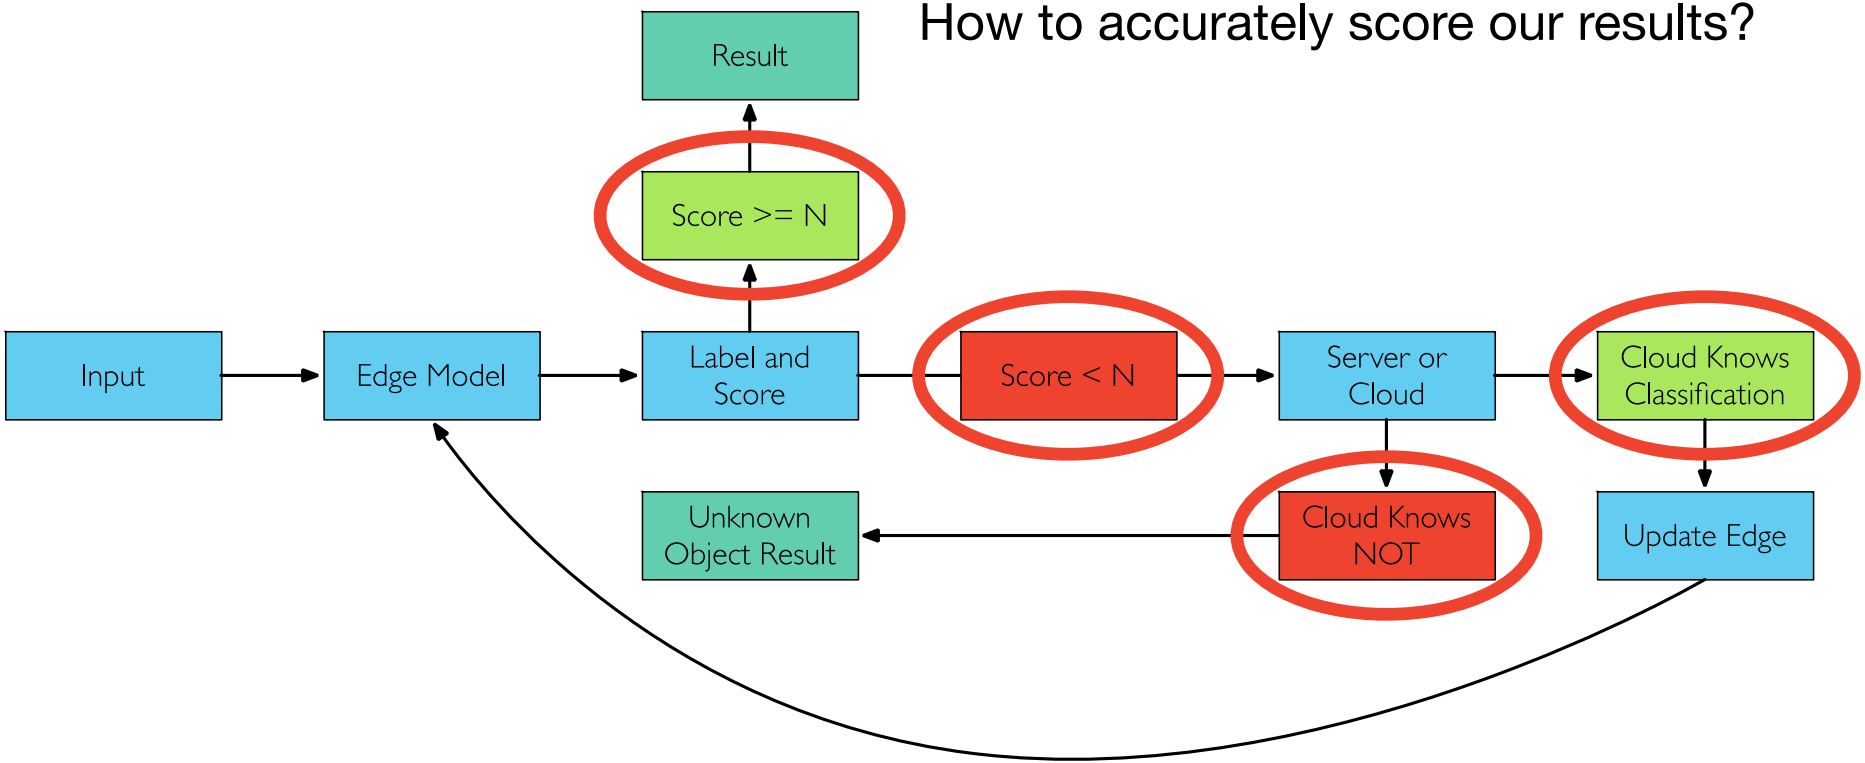

How to know you do not know?

- We need to know WHEN to go to the cloud for an update.
- Models tend to be overconfident in their results.
- Softmax is relative to KNOWN labels.
- Most objects probably UNKNOWN.
- model of all unknown labels…

Other Solutions?

- Model the Universe…

- …probably not practical on an embedded device yet.

Solutions – SVM Classifier

- Use the neural network as a feature extractor
- Reduces an image into a small vector
- The CNN output becomes the SVM input
- Measures how well the features fit “probability”

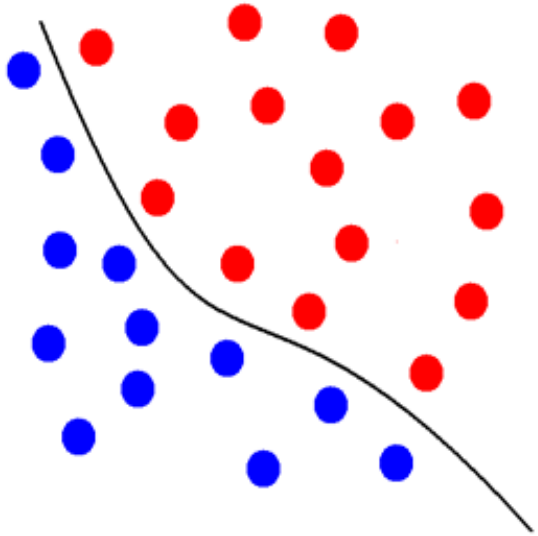

SVM Face Detection

- CNN can be used to extract features from an image
  - Trained to generate discriminating features
- SVM uses these features as inputs
  - Trained to fit a label and a **probability** from the input features
  - The probability is reliable and accurately reports unknown samples

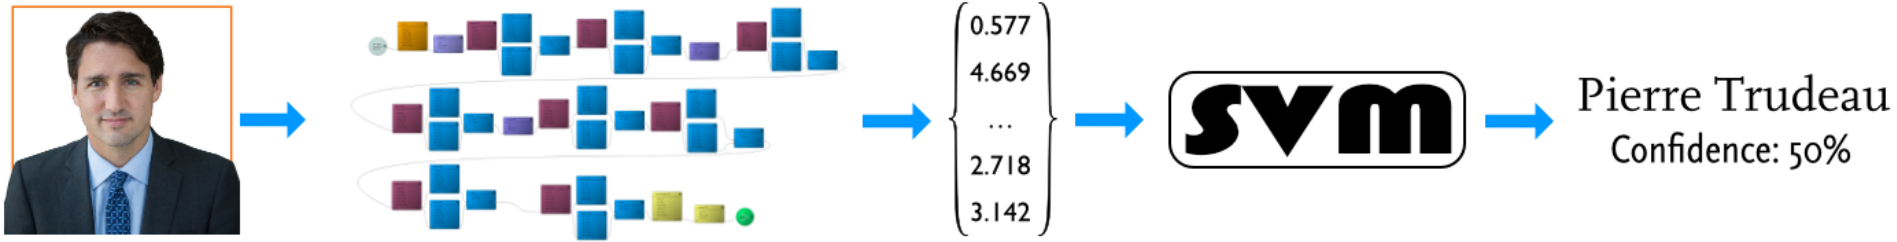

Distributing Model Updates

Architecture Flow – Model Update Size

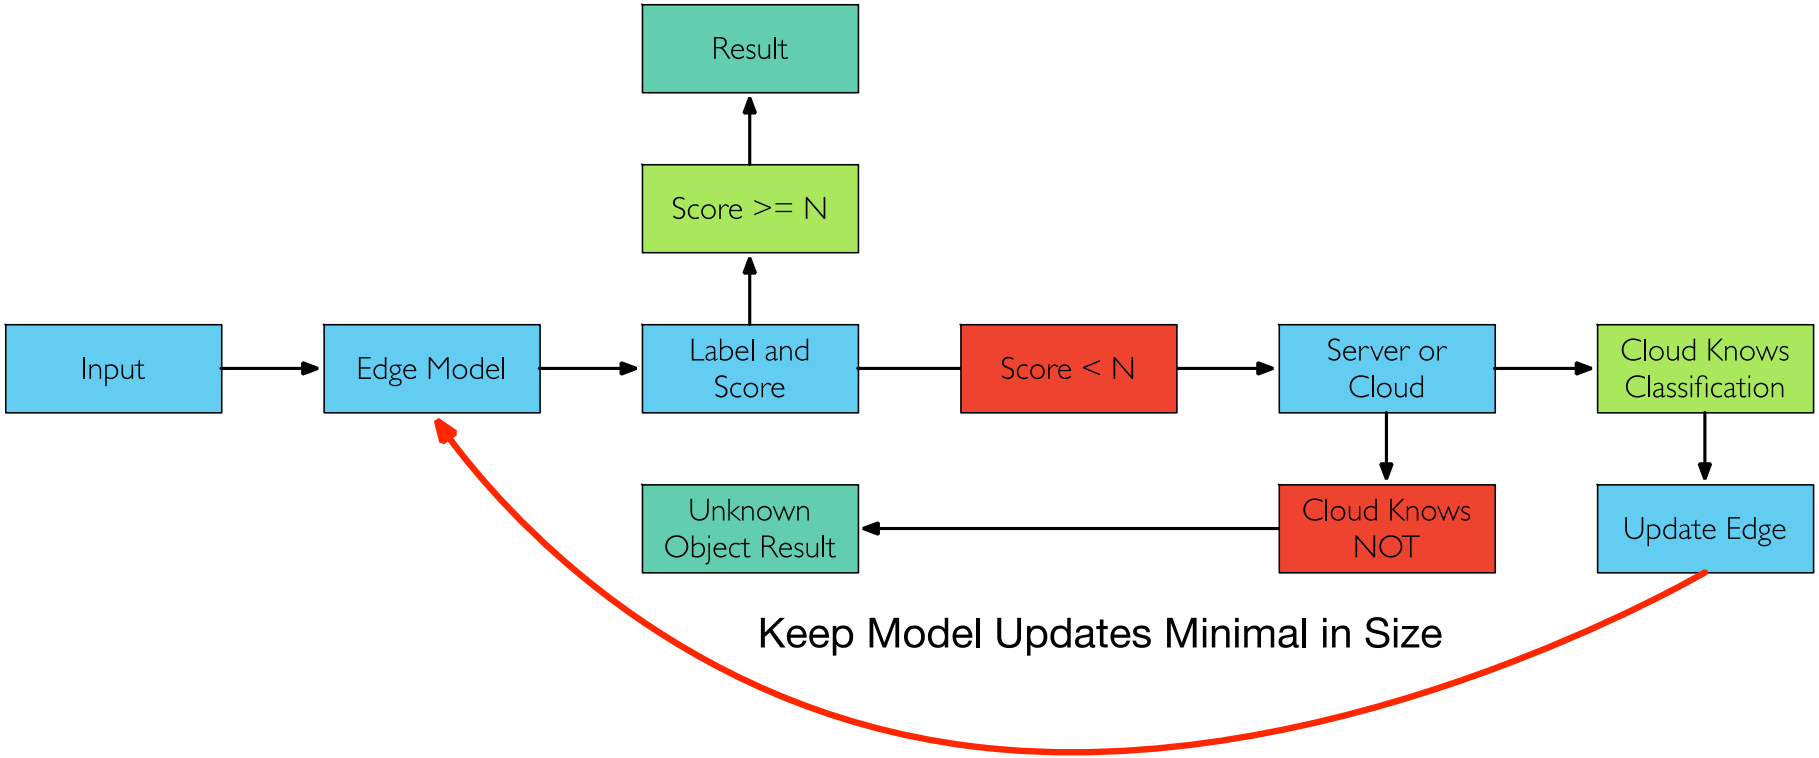

Updating the Edge

- IoT data transmission is still expensive
- Must Keep usage to a minimum

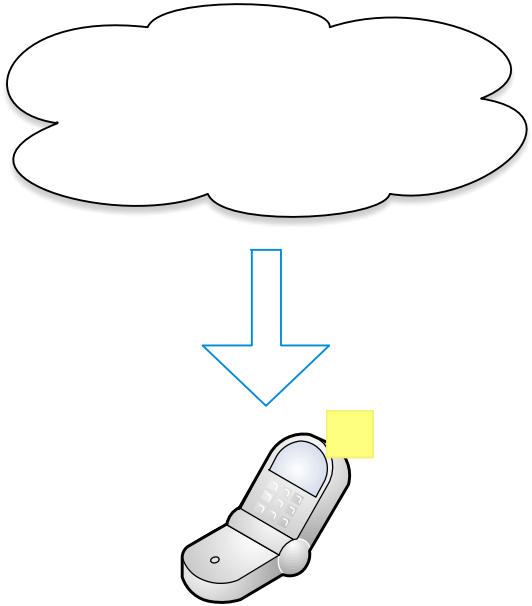

- Compress Models
  - Specifically compress the weights
- Send differences
  - Keep differences to a minimum

Limits of Compression

- Models generally do not compress well
  - Random data cannot be compressed
  - ”Smooth” data can be compressed
  - Are model weights more random or are they smooth?
- Lossy compression can greatly help
  - Neural networks are very resilient to error from lossy compression
  - Some models work well even down to 2 bits per weight!

Lossy Compression

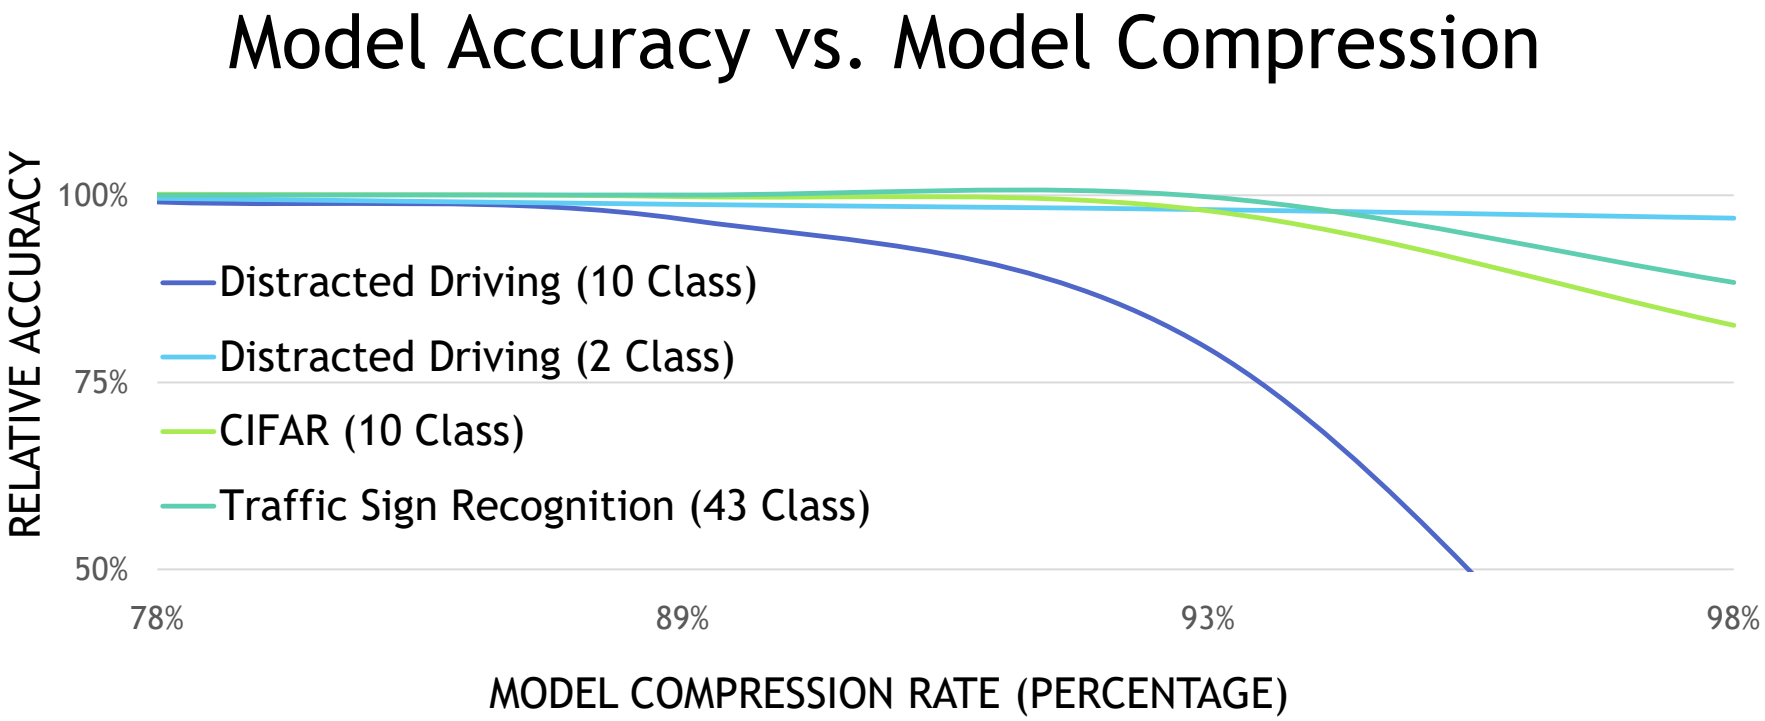

Transmit Differences

- Need to partially freeze model to avoid updating ALL weights
  - Cannot efficiently send differences if everything is changing
- If the front end is well generalized we only need to train the tail end
  - Same idea as transfer learning, in this case to help reduce data exchange
- In the SVM example we only need to update SVM weights, not the CNN model

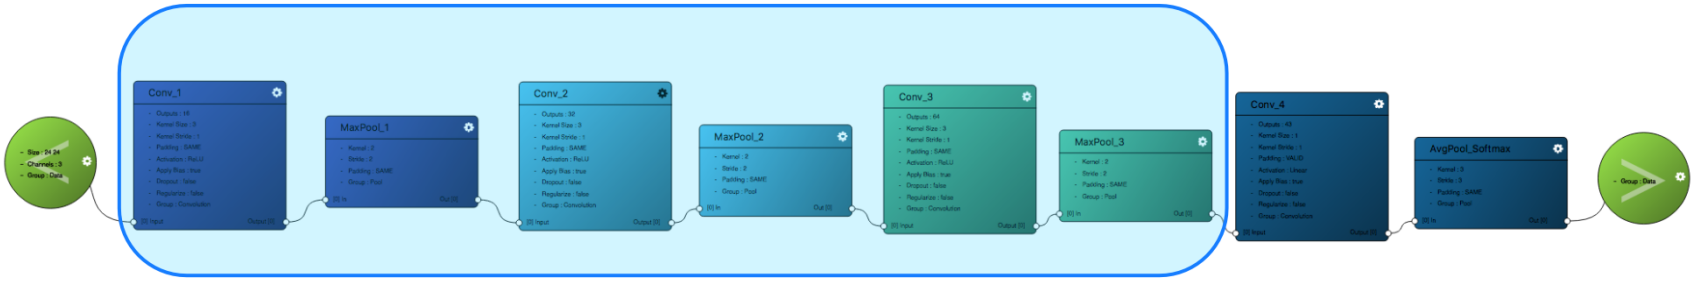

Frozen Model 23,472 Weights

Transfer 2816 Weights (10%)

Summary

Summary

- Cloud can use larger, evolving models
  - Can be used to train more focused models for the edge
  - Allows us to keep smaller models at the edge
- Need to know when to ask the cloud for help
  - Accurately detect when a sample is unknown
  - Go to the cloud for verification when unknown
  - Get updated models if available
- Need to efficiently distribute model updates
  - Lossy compression
  - Partial model updates

Resources

- www.au-zone.com
- www.embeddedml.com
- RT1050 https://www.youtube.com/watch?v=B2zwx6BYsKg
- i.MX8 Model Transfer https://www.youtube.com/watch?v=z0WtwXSlA9M
- DeepView MLTK https://www.youtube.com/watch?v=lS0QgM1VHaY
- Model Compression https://www.embedded-vision.com/platinum-members/embedded-vision-alliance/embedded-vision-training/documents/pages/deep-learning-software
- CNN Calibration https://arxiv.org/abs/1706.04599
- Modelling Uncertainty https://arxiv.org/abs/1509.05909